

Importing Libraries

In [1]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

Importing Datasets

In [2]:

```
df=pd.read_csv(r"C:\Users\user\Downloads\C10_air\csvs_per_year\csvs(Dataset)\madrid_2004.
df
```

Out[2]:

| | date | BEN | CO | EBE | MXY | NMHC | NO_2 | NOx | OXY | O_3 | |
|--------|---------------------|------|------|------|------|------|------------|------------|------|-----------|-----|
| 0 | 2004-08-01 01:00:00 | NaN | 0.66 | NaN | NaN | NaN | 89.550003 | 118.900002 | NaN | 40.020000 | 39 |
| 1 | 2004-08-01 01:00:00 | 2.66 | 0.54 | 2.99 | 6.08 | 0.18 | 51.799999 | 53.860001 | 3.28 | 51.689999 | 22 |
| 2 | 2004-08-01 01:00:00 | NaN | 1.02 | NaN | NaN | NaN | 93.389999 | 138.600006 | NaN | 20.860001 | 49 |
| 3 | 2004-08-01 01:00:00 | NaN | 0.53 | NaN | NaN | NaN | 87.290001 | 105.000000 | NaN | 36.730000 | 31 |
| 4 | 2004-08-01 01:00:00 | NaN | 0.17 | NaN | NaN | NaN | 34.910000 | 35.349998 | NaN | 86.269997 | 54 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 245491 | 2004-06-01 00:00:00 | 0.75 | 0.21 | 0.85 | 1.55 | 0.07 | 59.580002 | 64.389999 | 0.66 | 33.029999 | 30 |
| 245492 | 2004-06-01 00:00:00 | 2.49 | 0.75 | 2.44 | 4.57 | NaN | 97.139999 | 146.899994 | 2.34 | 7.740000 | 37 |
| 245493 | 2004-06-01 00:00:00 | NaN | NaN | NaN | NaN | 0.13 | 102.699997 | 132.600006 | NaN | 17.809999 | 22 |
| 245494 | 2004-06-01 00:00:00 | NaN | NaN | NaN | NaN | 0.09 | 82.599998 | 102.599998 | NaN | NaN | 45 |
| 245495 | 2004-06-01 00:00:00 | 3.01 | 0.67 | 2.78 | 5.12 | 0.20 | 92.550003 | 141.000000 | 2.60 | 11.460000 | 24 |

245496 rows × 17 columns



Data Cleaning and Data Preprocessing

In [3]:

```
df=df.dropna()
```

In [4]:

```
df.columns
```

Out[4]:

```
Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
      'PM10', 'PM25', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],
      dtype='object')
```

In [5]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 19397 entries, 5 to 245495
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        19397 non-null  object
1   BEN         19397 non-null  float64
2   CO          19397 non-null  float64
3   EBE         19397 non-null  float64
4   MXY         19397 non-null  float64
5   NMHC        19397 non-null  float64
6   NO_2        19397 non-null  float64
7   NOx         19397 non-null  float64
8   OXY         19397 non-null  float64
9   O_3         19397 non-null  float64
10  PM10        19397 non-null  float64
11  PM25        19397 non-null  float64
12  PXY         19397 non-null  float64
13  SO_2        19397 non-null  float64
14  TCH         19397 non-null  float64
15  TOL         19397 non-null  float64
16  station     19397 non-null  int64
dtypes: float64(15), int64(1), object(1)
memory usage: 2.7+ MB
```

In [8]:

```
data=df[['OXY', 'TCH', 'NOx']]
data
```

Out[8]:

| | OXY | TCH | NOx |
|--------|------|------|------------|
| 5 | 5.04 | 1.54 | 144.800003 |
| 22 | 0.50 | 1.55 | 32.799999 |
| 26 | 2.47 | 1.55 | 75.470001 |
| 32 | 2.56 | 1.67 | 165.800003 |
| 49 | 0.46 | 1.41 | 34.840000 |
| ... | ... | ... | ... |
| 245463 | 0.42 | 1.24 | 45.450001 |
| 245467 | 2.09 | 1.47 | 132.800003 |
| 245473 | 4.51 | 1.58 | 253.600006 |
| 245491 | 0.66 | 1.28 | 64.389999 |
| 245495 | 2.60 | 1.47 | 141.000000 |

19397 rows × 3 columns


Line chart

In [9]:

```
data.plot.line(subplots=True)
```

Out[9]:

array([<AxesSubplot:>, <AxesSubplot:>, <AxesSubplot:>], dtype=object)




Line chart

In [10]:

```
data.plot.line()
```

Out[10]:

<AxesSubplot:>



Bar chart

In [11]:


```
b=data[0:50]
```

In [12]:

```
b.plot.bar()
```

Out[12]:

<AxesSubplot:>




Histogram

In [13]:

```
data.plot.hist()
```

Out[13]:

<AxesSubplot:ylabel='Frequency'>




Area chart

In [14]:

```
data.plot.area()
```

Out[14]:

<AxesSubplot:>




Box chart

In [15]:

```
data.plot.box()
```

Out[15]:

<AxesSubplot:>




Pie chart

In [17]:

```
b.plot.pie(y='OXY' )
```

Out[17]:

<AxesSubplot:ylabel='OXY'>




Scatter chart

In [18]:

```
data.plot.scatter(x='OXY' ,y='NOx')
```

Out[18]:

<AxesSubplot:xlabel='OXY', ylabel='NOx'>



In [19]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 19397 entries, 5 to 245495
Data columns (total 17 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   date        19397 non-null  object  
 1   BEN         19397 non-null  float64 
 2   CO          19397 non-null  float64 
 3   EBE         19397 non-null  float64 
 4   MXY         19397 non-null  float64 
 5   NMHC        19397 non-null  float64 
 6   NO_2        19397 non-null  float64 
 7   NOx         19397 non-null  float64 
 8   OXY         19397 non-null  float64 
 9   O_3         19397 non-null  float64 
10  PM10        19397 non-null  float64 
11  PM25        19397 non-null  float64 
12  PXY         19397 non-null  float64 
13  SO_2        19397 non-null  float64 
14  TCH         19397 non-null  float64 
15  TOL         19397 non-null  float64 
16  station     19397 non-null  int64   
dtypes: float64(15), int64(1), object(1)
memory usage: 2.7+ MB
```

In [20]:

```
df.describe()
```

Out[20]:

| | BEN | CO | EBE | MXY | NMHC | NO_2 |
|-------|--------------|--------------|--------------|--------------|--------------|--------------|
| count | 19397.000000 | 19397.000000 | 19397.000000 | 19397.000000 | 19397.000000 | 19397.000000 |
| mean | 2.250781 | 0.675347 | 2.775913 | 5.424809 | 0.151024 | 62.887023 |
| std | 2.184724 | 0.591026 | 2.729622 | 5.554358 | 0.158603 | 37.952255 |
| min | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.090000 |
| 25% | 0.870000 | 0.320000 | 1.020000 | 1.780000 | 0.060000 | 35.150002 |
| 50% | 1.620000 | 0.520000 | 1.970000 | 3.800000 | 0.110000 | 58.310001 |
| 75% | 2.910000 | 0.860000 | 3.580000 | 7.260000 | 0.200000 | 85.730003 |
| max | 34.180000 | 8.900000 | 41.880001 | 91.599998 | 4.810000 | 355.100006 |

In [21]:

```
df1=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',  
        'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]
```

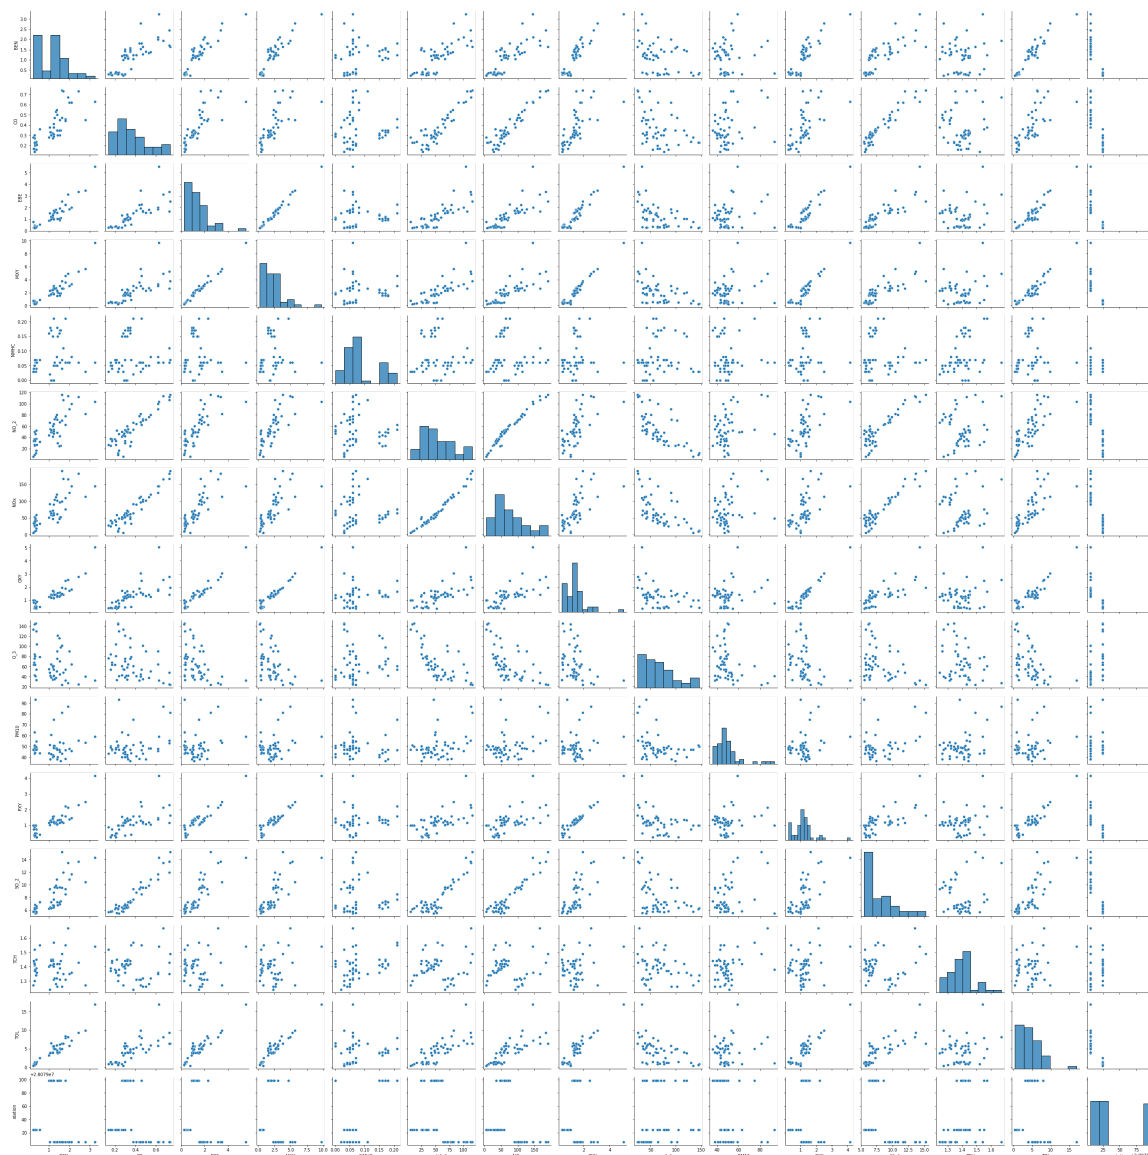
EDA AND VISUALIZATION

In [22]:

```
sns.pairplot(df1[0:50])
```

Out[22]:

<seaborn.axisgrid.PairGrid at 0x1f543b78100>



In [28]:

```
x=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',  
      'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]  
y=df['station']
```

In [29]:

```
from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

Linear Regression

In [30]:

```
from sklearn.linear_model import LinearRegression  
lr=LinearRegression()  
lr.fit(x_train,y_train)
```

Out[30]:

LinearRegression()

In [31]:

```
lr.intercept_
```

Out[31]:

28079074.466013305

In [32]:

```
coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
coeff
```

Out[32]:


| | Co-efficient |
|-------------|--------------|
| BEN | -3.428438 |
| CO | 23.720830 |
| EBE | 3.462263 |
| MXY | -3.428359 |
| NMHC | 82.207293 |
| NO_2 | -0.142861 |
| NOx | -0.245975 |
| OXY | -2.022878 |
| O_3 | -0.271436 |
| PM10 | 0.062757 |
| PXY | 6.011119 |
| SO_2 | -0.209234 |
| TCH | -6.418532 |
| TOL | 1.149801 |

In [33]:

```
prediction =lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[33]:

<matplotlib.collections.PathCollection at 0x1f550850700>



ACCURACY

In [34]:

```
lr.score(x_test,y_test)
```

Out[34]:

0.11870348140040887

In [35]:

```
lr.score(x_train,y_train)
```

Out[35]:

0.10067971554067745

Ridge and Lasso

In [36]:

```
from sklearn.linear_model import Ridge,Lasso
```

In [37]:

```
rr=Ridge(alpha=10)  
rr.fit(x_train,y_train)
```

Out[37]:

Ridge(alpha=10)

Accuracy(Ridge)

In [38]:

```
rr.score(x_test,y_test)
```

Out[38]:

0.11664945877101263

In [39]:

```
rr.score(x_train,y_train)
```

Out[39]:

0.1003371173292632

In [40]:

```
la=Lasso(alpha=10)  
la.fit(x_train,y_train)
```

Out[40]:

Lasso(alpha=10)

In [41]:

```
la.score(x_train,y_train)
```

Out[41]:

```
0.051781300470869596
```

Accuracy(Lasso)

In [42]:

```
la.score(x_test,y_test)
```

Out[42]:

```
0.056773795637802715
```

Accuracy(Elastic Net)

In [43]:

```
from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

Out[43]:

```
ElasticNet()
```

In [44]:

```
en.coef_
```

Out[44]:

```
array([-0.          ,  0.33534742,  1.45923881, -1.87938948,  0.          ,
        -0.1537371  , -0.09301587, -0.          , -0.20672745,  0.09436363,
         0.5749052  , -0.13540571,  0.          ,  1.20582375])
```

In [45]:

```
en.intercept_
```

Out[45]:

```
28079066.002937175
```

In [46]:

```
prediction=en.predict(x_test)
```


In [47]:

```
en.score(x_test,y_test)
```

Out[47]:

0.07033811305468773

Evaluation Metrics

In [48]:

```
from sklearn import metrics
print(metrics.mean_absolute_error(y_test,prediction))
print(metrics.mean_squared_error(y_test,prediction))
print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

38.45619093357262

1639.150228011821

40.4864202913992

Logistic Regression

In [49]:

```
from sklearn.linear_model import LogisticRegression
```

In [50]:

```
feature_matrix=df[['BEN', 'CO', 'EBE', 'MXV', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
                  'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
target_vector=df[ 'station']
```

In [51]:

```
feature_matrix.shape
```

Out[51]:

(19397, 14)

In [52]:

```
target_vector.shape
```

Out[52]:

(19397,)

In [53]:

```
from sklearn.preprocessing import StandardScaler
```

In [54]:

```
fs=StandardScaler().fit_transform(feature_matrix)
```

In [55]:

```
logr=LogisticRegression(max_iter=10000)  
logr.fit(fs,target_vector)
```

Out[55]:

```
LogisticRegression(max_iter=10000)
```

In [56]:

```
observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13,14]]
```

In [57]:

```
prediction=logr.predict(observation)  
print(prediction)
```

```
[28079006]
```

In [58]:

```
logr.classes_
```

Out[58]:

```
array([28079006, 28079024, 28079099], dtype=int64)
```

In [59]:

```
logr.score(fs,target_vector)
```

Out[59]:

```
0.7360416559261741
```

In [60]:

```
logr.predict_proba(observation)[0][0]
```

Out[60]:

```
0.9999978255573396
```

In [61]:

```
logr.predict_proba(observation)
```

Out[61]:

```
array([[9.99997826e-01, 7.75018107e-20, 2.17444266e-06]])
```

Random Forest

In [62]:

```
from sklearn.ensemble import RandomForestClassifier
```

In [63]:

```
rfc=RandomForestClassifier()  
rfc.fit(x_train,y_train)
```

Out[63]:

```
RandomForestClassifier()
```

In [64]:

```
parameters={'max_depth':[1,2,3,4,5],  
            'min_samples_leaf':[5,10,15,20,25],  
            'n_estimators':[10,20,30,40,50]  
}
```

In [65]:

```
from sklearn.model_selection import GridSearchCV  
grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")  
grid_search.fit(x_train,y_train)
```

Out[65]:

```
GridSearchCV(cv=2, estimator=RandomForestClassifier(),  
             param_grid={'max_depth': [1, 2, 3, 4, 5],  
                         'min_samples_leaf': [5, 10, 15, 20, 25],  
                         'n_estimators': [10, 20, 30, 40, 50]}},  
             scoring='accuracy')
```

In [66]:

```
grid_search.best_score_
```

Out[66]:

```
0.7720404784056986
```

In [67]:

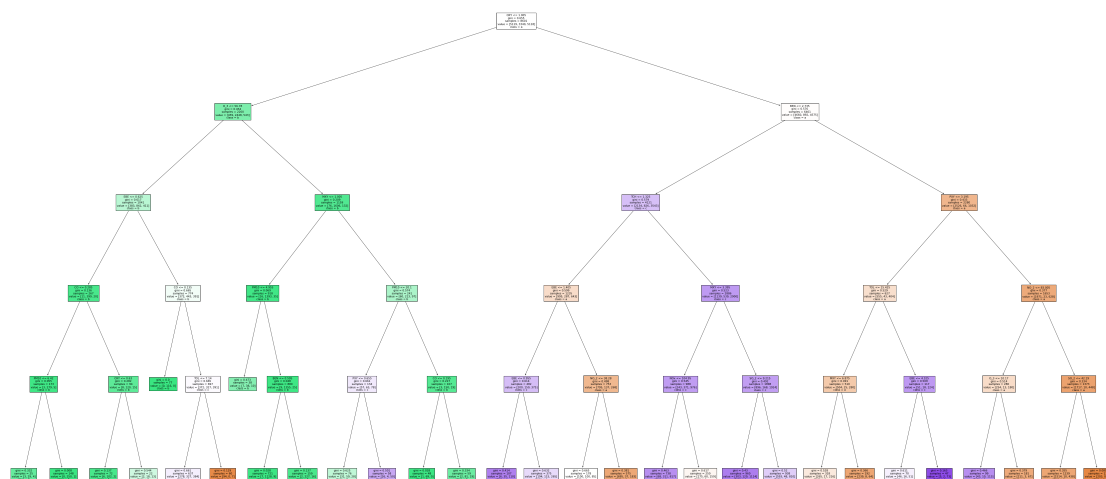
```
rfc_best=grid_search.best_estimator_
```

In [68]:

```

from sklearn.tree import plot_tree

plt.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['a','b','c','d'],f
39\nvalue = [1514, 10, 438]\nclass = a'),
Text(4384.285714285714, 181.19999999999982, 'gini = 0.089\nsamples = 13
4\nvalue = [203, 0, 10]\nclass = a')]
```



Conclusion

Accuracy

Linear Regression:0.10067971554067745

Ridge Regression:0.051781300470869596

Lasso Regression:0.056773795637802715

ElasticNet Regression:0.07033811305468773

Logistic Regression:0.7360416559261741

Random Forest:0.7720404784056986

Random Forest is suitable for this dataset