

Importing Libraries

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

Importing Datasets

```
In [2]: df=pd.read_csv(r"C:\Users\user\Downloads\C10_air\csvs_per_year\csvs(Dataset)\ma
df
```

Out[2]:

	date	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3	P
0	2004-08-01 01:00:00	NaN	0.66	NaN	NaN	NaN	89.550003	118.900002	NaN	40.020000	39.990000
1	2004-08-01 01:00:00	2.66	0.54	2.99	6.08	0.18	51.799999	53.860001	3.28	51.689999	22.950000
2	2004-08-01 01:00:00	NaN	1.02	NaN	NaN	NaN	93.389999	138.600006	NaN	20.860001	49.480000
3	2004-08-01 01:00:00	NaN	0.53	NaN	NaN	NaN	87.290001	105.000000	NaN	36.730000	31.070000
4	2004-08-01 01:00:00	NaN	0.17	NaN	NaN	NaN	34.910000	35.349998	NaN	86.269997	54.080000
...
245491	2004-06-01 00:00:00	0.75	0.21	0.85	1.55	0.07	59.580002	64.389999	0.66	33.029999	30.900000
245492	2004-06-01 00:00:00	2.49	0.75	2.44	4.57	NaN	97.139999	146.899994	2.34	7.740000	37.680000
245493	2004-06-01 00:00:00	NaN	NaN	NaN	NaN	0.13	102.699997	132.600006	NaN	17.809999	22.840000
245494	2004-06-01 00:00:00	NaN	NaN	NaN	NaN	0.09	82.599998	102.599998	NaN	NaN	45.630000
245495	2004-06-01 00:00:00	3.01	0.67	2.78	5.12	0.20	92.550003	141.000000	2.60	11.460000	24.380000

245496 rows × 17 columns



Data Cleaning and Data Preprocessing

```
In [3]: df=df.dropna()
```

```
In [4]: df.columns
```

```
Out[4]: Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
       'PM10', 'PM25', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],
      dtype='object')
```

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 19397 entries, 5 to 245495
Data columns (total 17 columns):
 #   Column    Non-Null Count  Dtype  
--- 
 0   date      19397 non-null   object 
 1   BEN        19397 non-null   float64
 2   CO         19397 non-null   float64
 3   EBE        19397 non-null   float64
 4   MXY        19397 non-null   float64
 5   NMHC       19397 non-null   float64
 6   NO_2       19397 non-null   float64
 7   NOx        19397 non-null   float64
 8   OXY        19397 non-null   float64
 9   O_3         19397 non-null   float64
 10  PM10       19397 non-null   float64
 11  PM25       19397 non-null   float64
 12  PXY        19397 non-null   float64
 13  SO_2       19397 non-null   float64
 14  TCH         19397 non-null   float64
 15  TOL         19397 non-null   float64
 16  station    19397 non-null   int64  
dtypes: float64(15), int64(1), object(1)
memory usage: 2.7+ MB
```

```
In [8]: data=df[['OXY', 'TCH', 'NOx']]  
data
```

Out[8]:

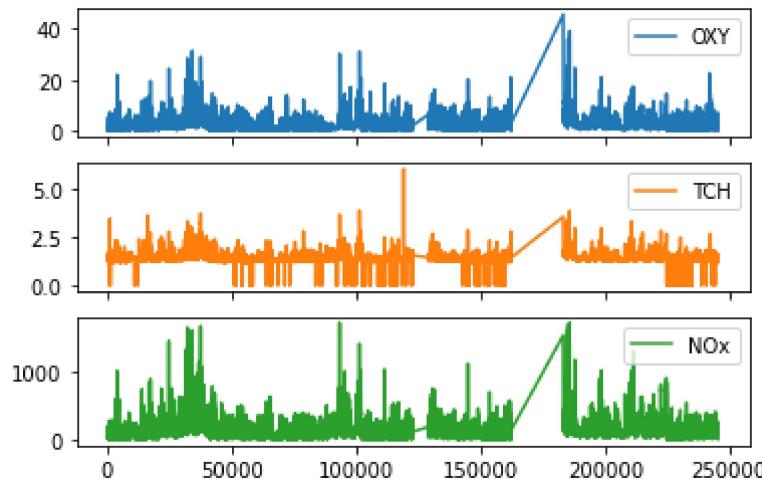
	OXY	TCH	NOx
5	5.04	1.54	144.800003
22	0.50	1.55	32.799999
26	2.47	1.55	75.470001
32	2.56	1.67	165.800003
49	0.46	1.41	34.840000
...
245463	0.42	1.24	45.450001
245467	2.09	1.47	132.800003
245473	4.51	1.58	253.600006
245491	0.66	1.28	64.389999
245495	2.60	1.47	141.000000

19397 rows × 3 columns

Line chart

```
In [9]: data.plot.line(subplots=True)
```

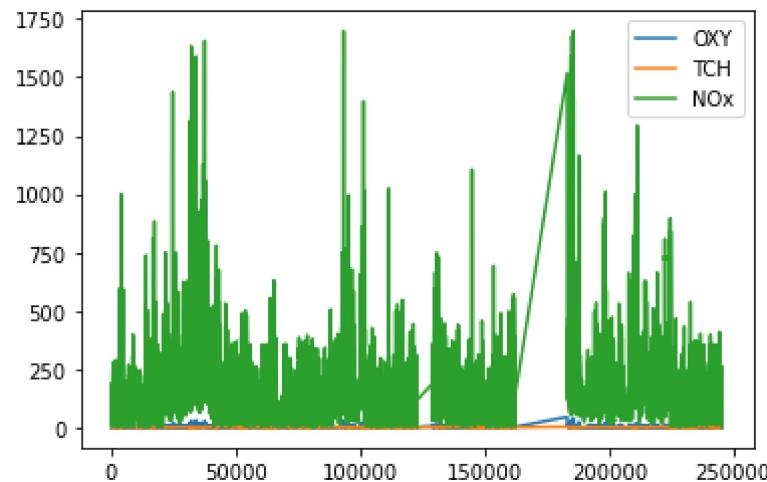
Out[9]: array([<AxesSubplot:>, <AxesSubplot:>, <AxesSubplot:>], dtype=object)



Line chart

```
In [10]: data.plot.line()
```

```
Out[10]: <AxesSubplot:>
```

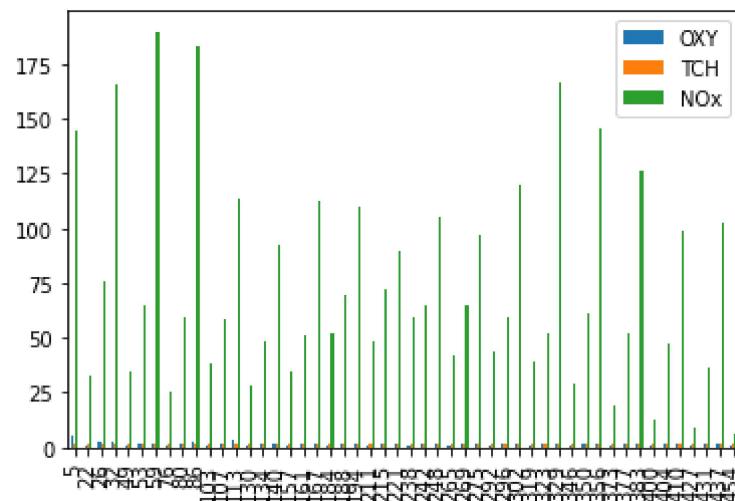


Bar chart

```
In [11]: b=data[0:50]
```

```
In [12]: b.plot.bar()
```

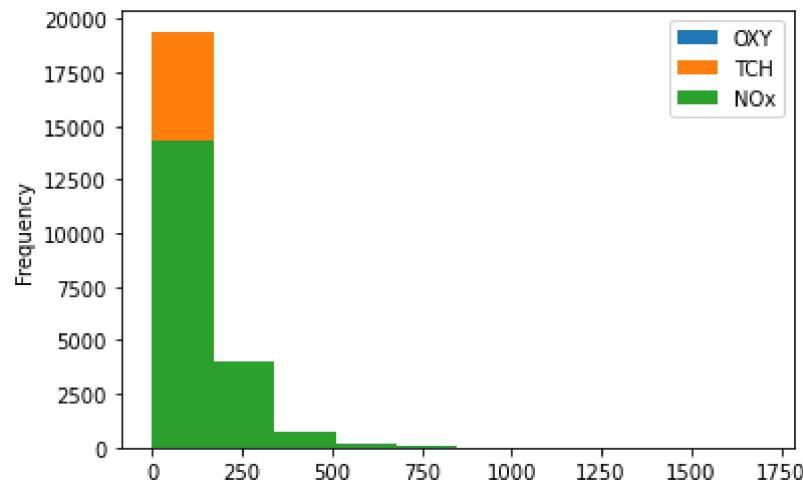
```
Out[12]: <AxesSubplot:>
```



Histogram

In [13]: `data.plot.hist()`

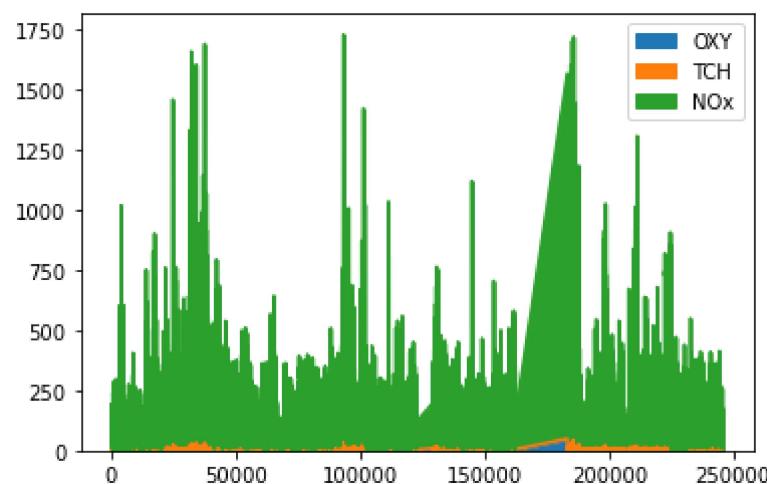
Out[13]: <AxesSubplot:ylabel='Frequency'>



Area chart

In [14]: `data.plot.area()`

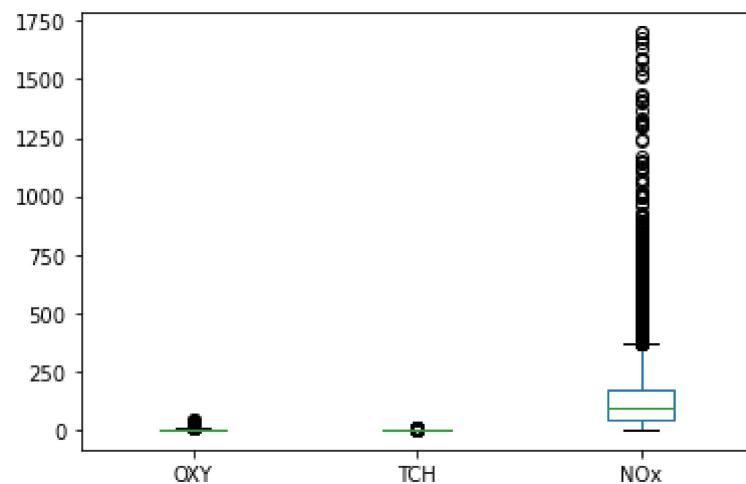
Out[14]: <AxesSubplot:>



Box chart

In [15]: `data.plot.box()`

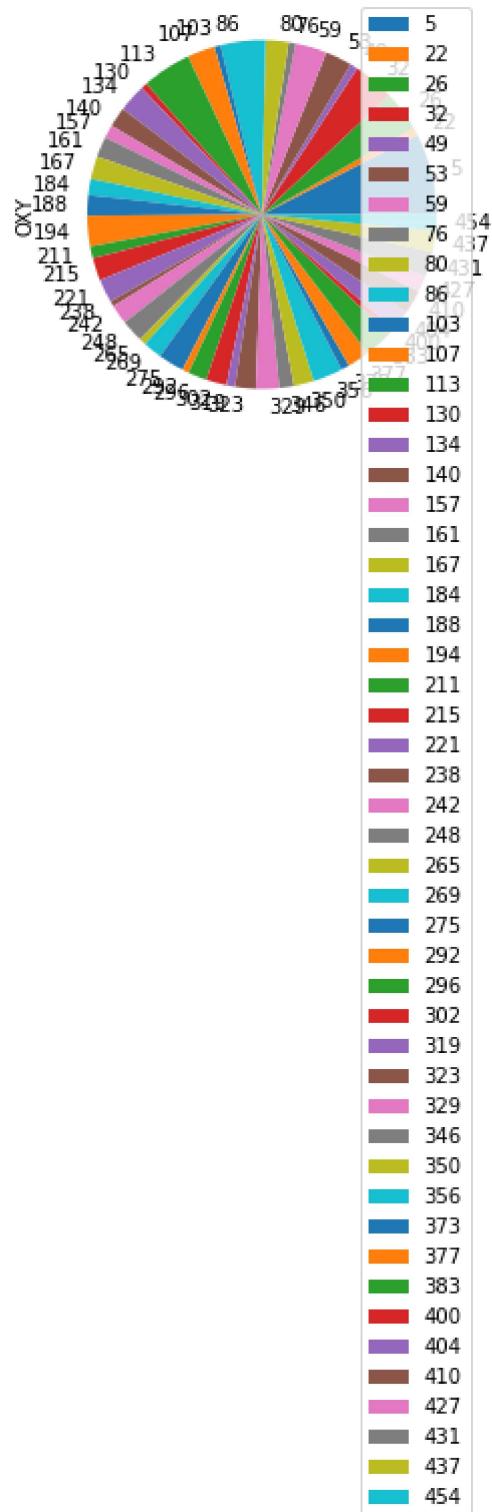
Out[15]: <AxesSubplot:>



Pie chart

```
In [17]: b.plot.pie(y='OXY' )
```

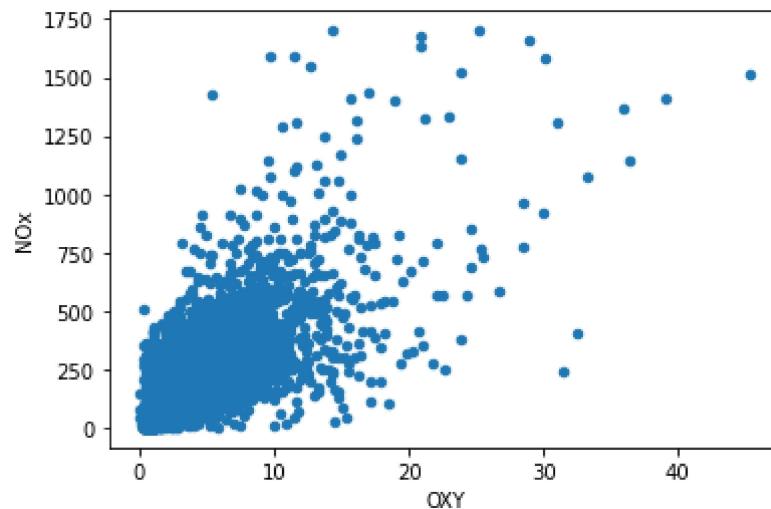
```
Out[17]: <AxesSubplot:ylabel='OXY'>
```



Scatter chart

```
In [18]: data.plot.scatter(x='OXY' ,y='NOx')
```

```
Out[18]: <AxesSubplot:xlabel='OXY', ylabel='NOx'>
```



```
In [19]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 19397 entries, 5 to 245495
Data columns (total 17 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   date        19397 non-null   object 
 1   BEN         19397 non-null   float64
 2   CO          19397 non-null   float64
 3   EBE         19397 non-null   float64
 4   MXY         19397 non-null   float64
 5   NMHC        19397 non-null   float64
 6   NO_2         19397 non-null   float64
 7   NOx         19397 non-null   float64
 8   OXY         19397 non-null   float64
 9   O_3          19397 non-null   float64
 10  PM10        19397 non-null   float64
 11  PM25        19397 non-null   float64
 12  PXY         19397 non-null   float64
 13  SO_2         19397 non-null   float64
 14  TCH          19397 non-null   float64
 15  TOL          19397 non-null   float64
 16  station      19397 non-null   int64  
dtypes: float64(15), int64(1), object(1)
memory usage: 2.7+ MB
```

In [20]: `df.describe()`

Out[20]:

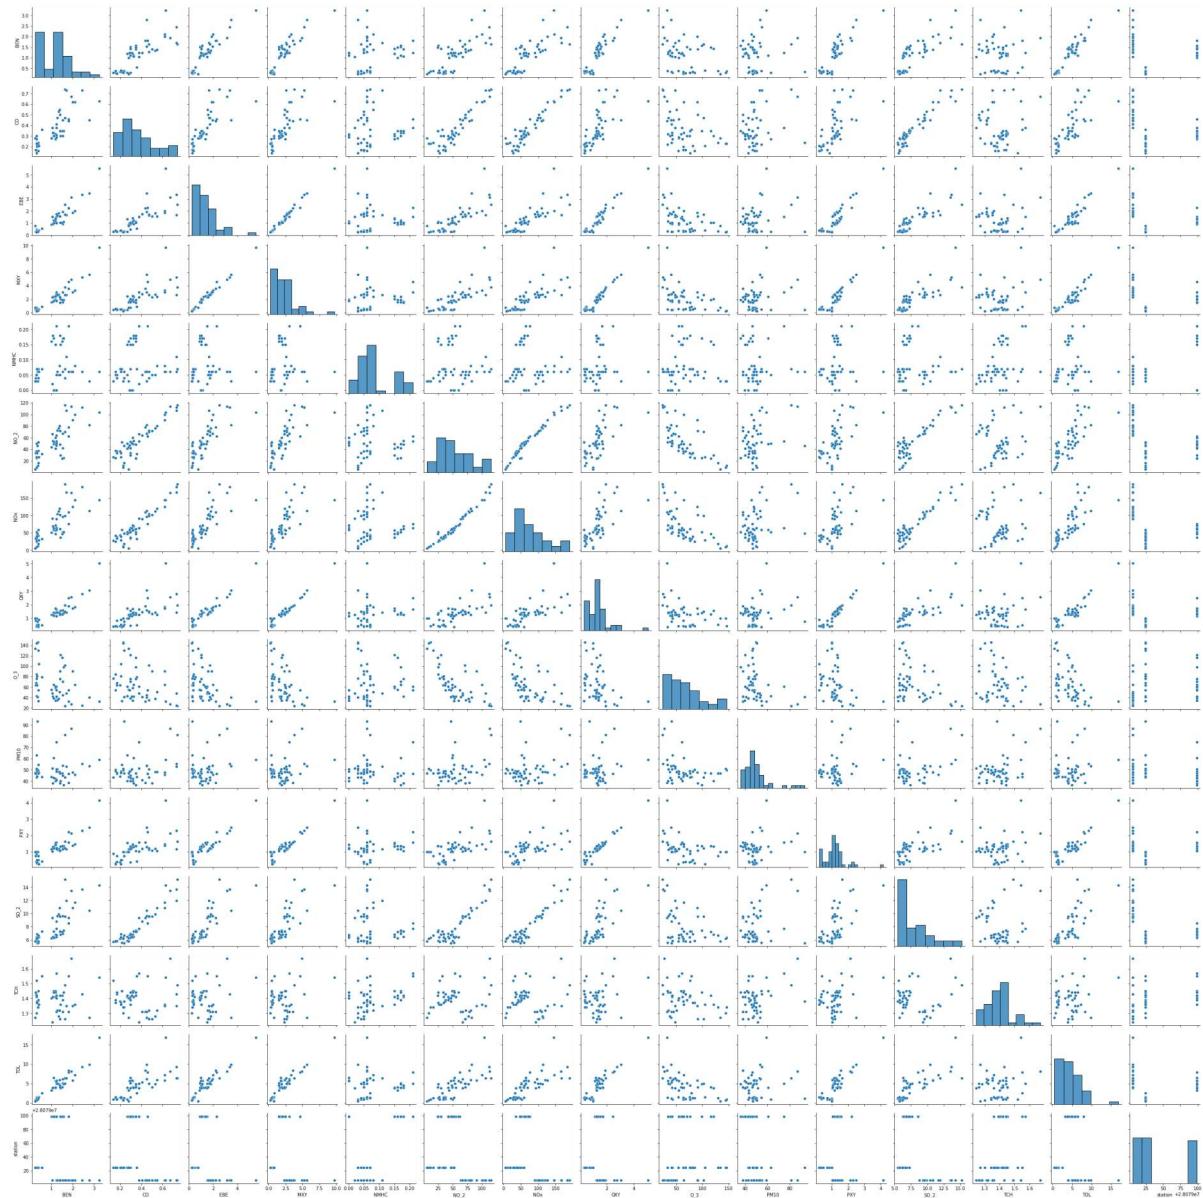
	BEN	CO	EBE	MXY	NMHC	NO_2	
count	19397.000000	19397.000000	19397.000000	19397.000000	19397.000000	19397.000000	193
mean	2.250781	0.675347	2.775913	5.424809	0.151024	62.887023	1
std	2.184724	0.591026	2.729622	5.554358	0.158603	37.952255	1
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.090000	
25%	0.870000	0.320000	1.020000	1.780000	0.060000	35.150002	
50%	1.620000	0.520000	1.970000	3.800000	0.110000	58.310001	
75%	2.910000	0.860000	3.580000	7.260000	0.200000	85.730003	1
max	34.180000	8.900000	41.880001	91.599998	4.810000	355.100006	17

In [21]: `df1=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3', 'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]`

EDA AND VISUALIZATION

```
In [22]: sns.pairplot(df1[0:50])
```

Out[22]: <seaborn.axisgrid.PairGrid at 0x1f543b78100>

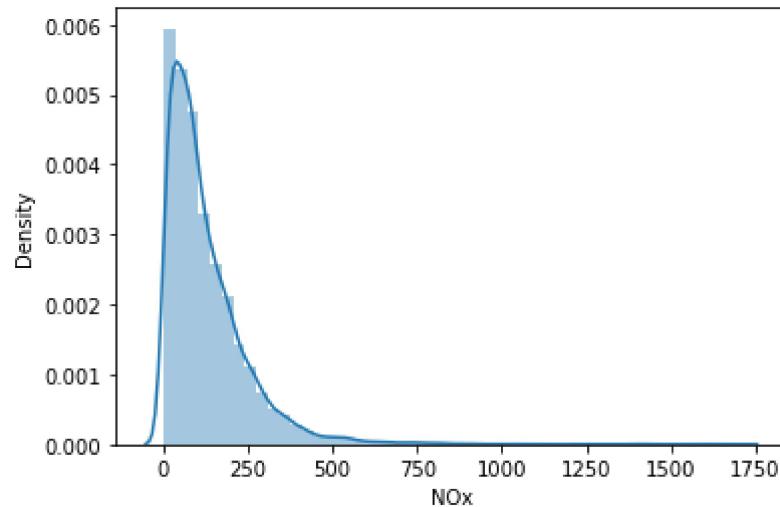


In [26]: `sns.distplot(df1['NOx'])`

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

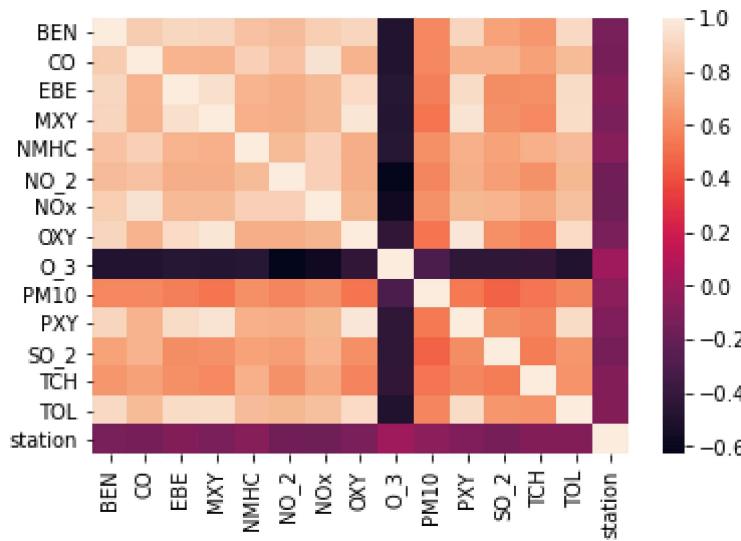
`warnings.warn(msg, FutureWarning)`

Out[26]: <AxesSubplot:xlabel='NOx', ylabel='Density'>



In [27]: `sns.heatmap(df1.corr())`

Out[27]: <AxesSubplot:>



TO TRAIN THE MODEL AND MODEL BUILDING

```
In [28]: x=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
           'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
y=df['station']
```

```
In [29]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

Linear Regression

```
In [30]: from sklearn.linear_model import LinearRegression
lr=LinearRegression()
lr.fit(x_train,y_train)
```

```
Out[30]: LinearRegression()
```

```
In [31]: lr.intercept_
```

```
Out[31]: 28079074.466013305
```

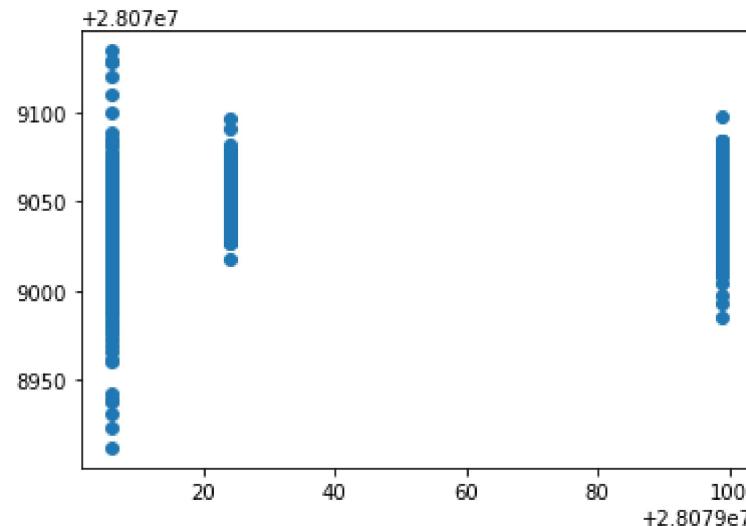
```
In [32]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
coeff
```

```
Out[32]:
```

	Co-efficient
BEN	-3.428438
CO	23.720830
EBE	3.462263
MXY	-3.428359
NMHC	82.207293
NO_2	-0.142861
NOx	-0.245975
OXY	-2.022878
O_3	-0.271436
PM10	0.062757
PXY	6.011119
SO_2	-0.209234
TCH	-6.418532
TOL	1.149801

```
In [33]: prediction = lr.predict(x_test)
plt.scatter(y_test,prediction)
```

```
Out[33]: <matplotlib.collections.PathCollection at 0x1f550850700>
```



ACCURACY

```
In [34]: lr.score(x_test,y_test)
```

```
Out[34]: 0.11870348140040887
```

```
In [35]: lr.score(x_train,y_train)
```

```
Out[35]: 0.10067971554067745
```

Ridge and Lasso

```
In [36]: from sklearn.linear_model import Ridge,Lasso
```

```
In [37]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

```
Out[37]: Ridge(alpha=10)
```

Accuracy(Ridge)

```
In [38]: rr.score(x_test,y_test)
```

```
Out[38]: 0.11664945877101263
```

```
In [39]: rr.score(x_train,y_train)
```

```
Out[39]: 0.1003371173292632
```

```
In [40]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

```
Out[40]: Lasso(alpha=10)
```

```
In [41]: la.score(x_train,y_train)
```

```
Out[41]: 0.051781300470869596
```

Accuracy(Lasso)

```
In [42]: la.score(x_test,y_test)
```

```
Out[42]: 0.056773795637802715
```

Accuracy(Elastic Net)

```
In [43]: from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

```
Out[43]: ElasticNet()
```

```
In [44]: en.coef_
```

```
Out[44]: array([-0.          ,  0.33534742,  1.45923881, -1.87938948,  0.          ,
 -0.1537371 , -0.09301587, -0.          , -0.20672745,  0.09436363,
 0.5749052 , -0.13540571,  0.          ,  1.20582375])
```

```
In [45]: en.intercept_
```

```
Out[45]: 28079066.002937175
```

```
In [46]: prediction=en.predict(x_test)
```

```
In [47]: en.score(x_test,y_test)
```

```
Out[47]: 0.07033811305468773
```

Evaluation Metrics

```
In [48]: from sklearn import metrics
print(metrics.mean_absolute_error(y_test,prediction))
print(metrics.mean_squared_error(y_test,prediction))
print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

38.45619093357262
1639.150228011821
40.4864202913992

Logistic Regression

```
In [49]: from sklearn.linear_model import LogisticRegression
```

```
In [50]: feature_matrix=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_P10', 'PXY', 'SO_2', 'TCH', 'TOL']]
target_vector=df[ 'station']
```

```
In [51]: feature_matrix.shape
```

```
Out[51]: (19397, 14)
```

```
In [52]: target_vector.shape
```

```
Out[52]: (19397,)
```

```
In [53]: from sklearn.preprocessing import StandardScaler
```

```
In [54]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [55]: logr=LogisticRegression(max_iter=10000)
logr.fit(fs,target_vector)
```

```
Out[55]: LogisticRegression(max_iter=10000)
```

```
In [56]: observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13,14]]
```

```
In [57]: prediction=logr.predict(observation)
print(prediction)
```

[28079006]

```
In [58]: logr.classes_
```

```
Out[58]: array([28079006, 28079024, 28079099], dtype=int64)
```

```
In [59]: logr.score(fs,target_vector)
```

```
Out[59]: 0.7360416559261741
```

```
In [60]: logr.predict_proba(observation)[0][0]
```

```
Out[60]: 0.9999978255573396
```

```
In [61]: logr.predict_proba(observation)
```

```
Out[61]: array([[9.99997826e-01, 7.75018107e-20, 2.17444266e-06]])
```

Random Forest

```
In [62]: from sklearn.ensemble import RandomForestClassifier
```

```
In [63]: rfc=RandomForestClassifier()  
rfc.fit(x_train,y_train)
```

```
Out[63]: RandomForestClassifier()
```

```
In [64]: parameters={'max_depth':[1,2,3,4,5],  
                  'min_samples_leaf':[5,10,15,20,25],  
                  'n_estimators':[10,20,30,40,50]  
}
```

```
In [65]: from sklearn.model_selection import GridSearchCV  
grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")  
grid_search.fit(x_train,y_train)
```

```
Out[65]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),  
                      param_grid={'max_depth': [1, 2, 3, 4, 5],  
                                  'min_samples_leaf': [5, 10, 15, 20, 25],  
                                  'n_estimators': [10, 20, 30, 40, 50]},  
                      scoring='accuracy')
```

```
In [66]: grid_search.best_score_
```

```
Out[66]: 0.7720404784056986
```

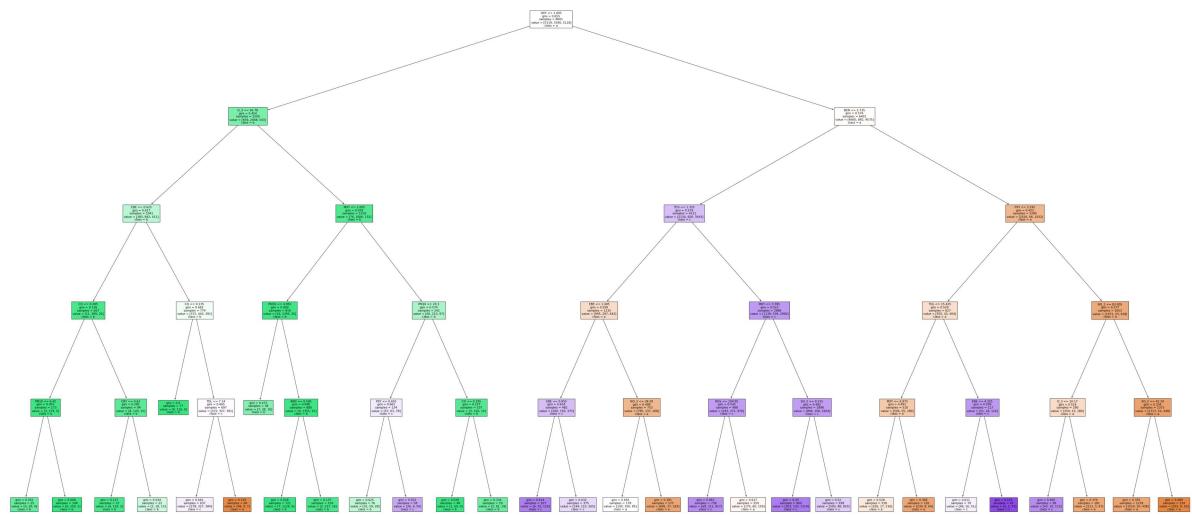
```
In [67]: rfc_best=grid_search.best_estimator_
```

```
In [68]: from sklearn.tree import plot_tree  
  
plt.figure(figsize=(80,40))  
plot_tree(rfc_best.estimators_[5], feature_names=x.columns, class_names=['a', 'b'])
```

```
Out[68]: [Text(2052.642857142857, 1993.2, 'OXY <= 1.005\ngini = 0.655\nsamples = 8601\nvalue = [5119, 3340, 5118]\nclass = a'),  
 Text(916.7142857142857, 1630.8000000000002, 'O_3 <= 56.78\ngini = 0.454\nsamples = 2200\nvalue = [459, 2448, 543]\nclass = b'),  
 Text(518.1428571428571, 1268.4, 'EBE <= 0.625\ngini = 0.617\nsamples = 1041\nvalue = [383, 842, 411]\nclass = b'),  
 Text(318.85714285714283, 906.0, 'CO <= 0.285\ngini = 0.136\nsamples = 267\nvalue = [11, 399, 20]\nclass = b'),  
 Text(159.42857142857142, 543.5999999999999, 'PM10 <= 6.42\ngini = 0.055\nsamples = 173\nvalue = [3, 279, 5]\nclass = b'),  
 Text(79.71428571428571, 181.1999999999982, 'gini = 0.332\nsamples = 25\nvalue = [3, 29, 4]\nclass = b'),  
 Text(239.1428571428571, 181.1999999999982, 'gini = 0.008\nsamples = 148\nvalue = [0, 250, 1]\nclass = b'),  
 Text(478.2857142857142, 543.5999999999999, 'OXY <= 0.63\ngini = 0.282\nsamples = 94\nvalue = [8, 120, 15]\nclass = b'),  
 Text(398.57142857142856, 181.1999999999982, 'gini = 0.137\nsamples = 72\nvalue = [6, 102, 2]\nclass = b'),  
 Text(558.0, 181.1999999999982, 'gini = 0.544\nsamples = 22\nvalue = [2, 18, 13]\nclass = b'),  
 Text(717.4285714285713, 906.0, 'CO <= 0.135\ngini = 0.665\nsamples = 774\nvalue = [372, 443, 391]\nclass = b'),  
 Text(637.7142857142857, 543.5999999999999, 'gini = 0.0\nsamples = 77\nvalue = [0, 116, 0]\nclass = b'),  
 Text(797.1428571428571, 543.5999999999999, 'TOL <= 7.14\ngini = 0.665\nsamples = 697\nvalue = [372, 327, 391]\nclass = c'),  
 Text(717.4285714285713, 181.1999999999982, 'gini = 0.661\nsamples = 637\nvalue = [278, 327, 384]\nclass = c'),  
 Text(876.8571428571428, 181.1999999999982, 'gini = 0.129\nsamples = 60\nvalue = [94, 0, 7]\nclass = a'),  
 Text(1315.2857142857142, 1268.4, 'MXY <= 1.005\ngini = 0.209\nsamples = 1159\nvalue = [76, 1606, 132]\nclass = b'),  
 Text(1036.2857142857142, 906.0, 'PM10 <= 4.955\ngini = 0.069\nsamples = 918\nvalue = [16, 1393, 35]\nclass = b'),  
 Text(956.5714285714284, 543.5999999999999, 'gini = 0.473\nsamples = 38\nvalue = [7, 38, 10]\nclass = b'),  
 Text(1116.0, 543.5999999999999, 'BEN <= 0.595\ngini = 0.048\nsamples = 880\nvalue = [9, 1355, 25]\nclass = b'),  
 Text(1036.2857142857142, 181.1999999999982, 'gini = 0.028\nsamples = 721\nvalue = [7, 1128, 9]\nclass = b'),  
 Text(1195.7142857142856, 181.1999999999982, 'gini = 0.137\nsamples = 159\nvalue = [2, 227, 16]\nclass = b'),  
 Text(1594.2857142857142, 906.0, 'PM10 <= 20.1\ngini = 0.574\nsamples = 241\nvalue = [60, 213, 97]\nclass = b'),  
 Text(1434.8571428571427, 543.5999999999999, 'PXY <= 0.655\ngini = 0.661\nsamples = 134\nvalue = [57, 63, 78]\nclass = c'),  
 Text(1355.142857142857, 181.1999999999982, 'gini = 0.625\nsamples = 76\nvalue = [31, 59, 28]\nclass = b'),  
 Text(1514.5714285714284, 181.1999999999982, 'gini = 0.501\nsamples = 58\nvalue = [26, 4, 50]\nclass = c'),  
 Text(1753.7142857142856, 543.5999999999999, 'CO <= 0.195\ngini = 0.227\nsamples = 107\nvalue = [3, 150, 19]\nclass = b'),  
 Text(1673.9999999999998, 181.1999999999982, 'gini = 0.028\nsamples = 48\nvalue = [1, 69, 0]\nclass = b'),  
 Text(1833.4285714285713, 181.1999999999982, 'gini = 0.334\nsamples = 59\nvalue = [2, 81, 19]\nclass = b'),  
 Text(3188.5714285714284, 1630.8000000000002, 'BEN <= 2.735\ngini = 0.576\nsa
```

```
mples = 6401\nvalue = [4660, 892, 4575]\nclass = a'),  
    Text(2550.8571428571427, 1268.4, 'TCH <= 1.325\ngini = 0.579\nsamples = 4121  
\nvalue = [2134, 826, 3543]\nclass = c'),  
    Text(2232.0, 906.0, 'EBE <= 1.405\ngini = 0.599\nsamples = 1235\nvalue = [99  
5, 287, 643]\nclass = a'),  
    Text(2072.5714285714284, 543.5999999999999, 'EBE <= 0.955\ngini = 0.614\nsam  
ples = 482\nvalue = [200, 150, 375]\nclass = c'),  
    Text(1992.8571428571427, 181.1999999999982, 'gini = 0.414\nsamples = 107\nv  
alue = [6, 35, 110]\nclass = c'),  
    Text(2152.285714285714, 181.1999999999982, 'gini = 0.632\nsamples = 375\nv  
alue = [194, 115, 265]\nclass = c'),  
    Text(2391.428571428571, 543.5999999999999, 'NO_2 <= 38.28\ngini = 0.498\nsam  
ples = 753\nvalue = [795, 137, 268]\nclass = a'),  
    Text(2311.7142857142853, 181.1999999999982, 'gini = 0.665\nsamples = 178\nv  
alue = [100, 100, 85]\nclass = a'),  
    Text(2471.142857142857, 181.1999999999982, 'gini = 0.381\nsamples = 575\nv  
alue = [695, 37, 183]\nclass = a'),  
    Text(2869.7142857142853, 906.0, 'MXY <= 3.395\ngini = 0.523\nsamples = 2886  
\nvalue = [1139, 539, 2900]\nclass = c'),  
    Text(2710.285714285714, 543.5999999999999, 'NOx <= 104.95\ngini = 0.545\nsam  
ples = 988\nvalue = [243, 371, 976]\nclass = c'),  
    Text(2630.5714285714284, 181.1999999999982, 'gini = 0.463\nsamples = 738\nv  
alue = [68, 311, 817]\nclass = c'),  
    Text(2790.0, 181.1999999999982, 'gini = 0.617\nsamples = 250\nvalue = [175,  
60, 159]\nclass = a'),  
    Text(3029.142857142857, 543.5999999999999, 'SO_2 <= 9.215\ngini = 0.492\nsam  
ples = 1898\nvalue = [896, 168, 1924]\nclass = c'),  
    Text(2949.428571428571, 181.1999999999982, 'gini = 0.43\nsamples = 960\nv  
alue = [303, 120, 1114]\nclass = c'),  
    Text(3108.8571428571427, 181.1999999999982, 'gini = 0.52\nsamples = 938\nv  
alue = [593, 48, 810]\nclass = c'),  
    Text(3826.2857142857138, 1268.4, 'PXY <= 3.195\ngini = 0.433\nsamples = 2280  
\nvalue = [2526, 66, 1032]\nclass = a'),  
    Text(3507.428571428571, 906.0, 'TOL <= 15.425\ngini = 0.529\nsamples = 627\nn  
value = [555, 43, 404]\nclass = a'),  
    Text(3347.999999999995, 543.5999999999999, 'MXY <= 6.875\ngini = 0.491\nsam  
ples = 510\nvalue = [504, 25, 280]\nclass = a'),  
    Text(3268.285714285714, 181.1999999999982, 'gini = 0.528\nsamples = 318\nv  
alue = [265, 17, 216]\nclass = a'),  
    Text(3427.7142857142853, 181.1999999999982, 'gini = 0.366\nsamples = 192\nv  
alue = [239, 8, 64]\nclass = a'),  
    Text(3666.8571428571427, 543.5999999999999, 'EBE <= 4.325\ngini = 0.509\nsam  
ples = 117\nvalue = [51, 18, 124]\nclass = c'),  
    Text(3587.142857142857, 181.1999999999982, 'gini = 0.611\nsamples = 70\nv  
alue = [46, 16, 51]\nclass = c'),  
    Text(3746.5714285714284, 181.1999999999982, 'gini = 0.163\nsamples = 47\nv  
alue = [5, 2, 73]\nclass = c'),  
    Text(4145.142857142857, 906.0, 'NO_2 <= 83.005\ngini = 0.377\nsamples = 1653  
\nvalue = [1971, 23, 628]\nclass = a'),  
    Text(3985.7142857142853, 543.5999999999999, 'O_3 <= 10.17\ngini = 0.514\nsam  
ples = 280\nvalue = [254, 13, 180]\nclass = a'),  
    Text(3905.999999999995, 181.1999999999982, 'gini = 0.466\nsamples = 99\nv  
alue = [43, 10, 113]\nclass = c'),  
    Text(4065.428571428571, 181.1999999999982, 'gini = 0.379\nsamples = 181\nv  
alue = [211, 3, 67]\nclass = a'),  
    Text(4304.571428571428, 543.5999999999999, 'SO_2 <= 42.19\ngini = 0.334\nsam  
ples = 1373\nvalue = [1717, 10, 448]\nclass = a'),
```

```
Text(4224.857142857142, 181.1999999999982, 'gini = 0.355\nsamples = 1239\nvalue = [1514, 10, 438]\nnclass = a'),  
Text(4384.285714285714, 181.1999999999982, 'gini = 0.089\nsamples = 134\nvalue = [203, 0, 10]\nnclass = a')]
```



Conclusion

Accuracy

Linear Regression:0.10067971554067745

Ridge Regression:0.051781300470869596

Lasso Regression:0.056773795637802715

ElasticNet Regression:0.07033811305468773

Logistic Regression:0.7360416559261741

Random Forest:0.7720404784056986

Random Forest is suitable for this dataset