Importing Libraries

In [1]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

Importing Datasets

In [2]:

```
df=pd.read_csv(r"C:\Users\user\Downloads\C10_air\csvs_per_year\csvs(Dataset)\madrid_2006.
df
```

Out[2]:

	date	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3	
0	2006-02-01 01:00:00	NaN	1.84	NaN	NaN	NaN	155.100006	490.100006	NaN	4.880000	97
1	2006-02-01 01:00:00	1.68	1.01	2.38	6.36	0.32	94.339996	229.699997	3.04	7.100000	25
2	2006-02-01 01:00:00	NaN	1.25	NaN	NaN	NaN	66.800003	192.000000	NaN	4.430000	34
3	2006-02-01 01:00:00	NaN	1.68	NaN	NaN	NaN	103.000000	407.799988	NaN	4.830000	28
4	2006-02-01 01:00:00	NaN	1.31	NaN	NaN	NaN	105.400002	269.200012	NaN	6.990000	54
...	...	...	...	...	...	...	...	...	...	...	
230563	2006-05-01 00:00:00	5.88	0.83	6.23	NaN	0.20	112.500000	218.000000	NaN	24.389999	93
230564	2006-05-01 00:00:00	0.76	0.32	0.48	1.09	0.08	51.900002	54.820000	0.61	48.410000	29
230565	2006-05-01 00:00:00	0.96	NaN	0.69	NaN	0.19	135.100006	179.199997	NaN	11.460000	64
230566	2006-05-01 00:00:00	0.50	NaN	0.67	NaN	0.10	82.599998	105.599998	NaN	NaN	94
230567	2006-05-01 00:00:00	1.95	0.74	1.99	4.00	0.24	107.300003	160.199997	2.01	17.730000	52

230568 rows × 17 columns

Data Cleaning and Data Preprocessing

In [3]:

```
df=df.dropna()
```

In [4]:

```
df.columns
```

Out[4]:

```
Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O
_3',
       'PM10', 'PM25', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],
      dtype='object')
```

In [5]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 24758 entries, 5 to 230567
Data columns (total 17 columns):
 #   Column   Non-Null Count  Dtype  
---  ------   --------------  -----  
 0   date     24758 non-null  object 
 1   BEN      24758 non-null  float64
 2   CO       24758 non-null  float64
 3   EBE      24758 non-null  float64
 4   MXY      24758 non-null  float64
 5   NMHC     24758 non-null  float64
 6   NO_2     24758 non-null  float64
 7   NOx      24758 non-null  float64
 8   OXY      24758 non-null  float64
 9   O_3      24758 non-null  float64
 10  PM10     24758 non-null  float64
 11  PM25     24758 non-null  float64
 12  PXY      24758 non-null  float64
 13  SO_2     24758 non-null  float64
 14  TCH      24758 non-null  float64
 15  TOL      24758 non-null  float64
 16  station  24758 non-null  int64  
dtypes: float64(15), int64(1), object(1)
memory usage: 3.4+ MB
```

In [7]:

```
data=df[['NMHC', 'NO_2', 'O_3']]
data
```

Out[7]:

	NMHC	NO_2	O_3
5	0.44	142.199997	5.990000
22	0.17	59.910000	2.450000
25	0.40	117.699997	4.780000
31	0.25	92.059998	5.920000
48	0.16	60.189999	2.280000
...	...	...	...
230538	0.10	49.259998	64.599998
230541	0.33	63.220001	17.670000
230547	0.26	202.399994	11.130000
230564	0.08	51.900002	48.410000
230567	0.24	107.300003	17.730000

24758 rows × 3 columns

Line chart

In [8]:

```
data.plot.line(subplots=True)
```

Out[8]:

```
array([<AxesSubplot:>, <AxesSubplot:>, <AxesSubplot:>], dtype=object)
```

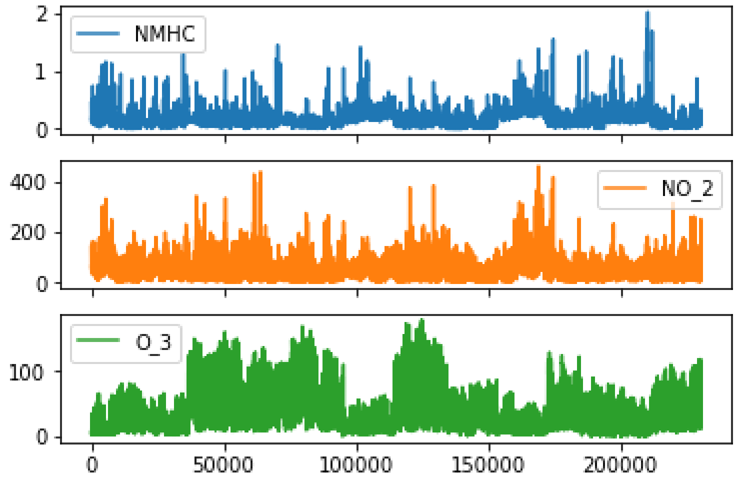

Line chart

In [10]:

```
data.plot.line()
```

Out[10]:

<AxesSubplot:>

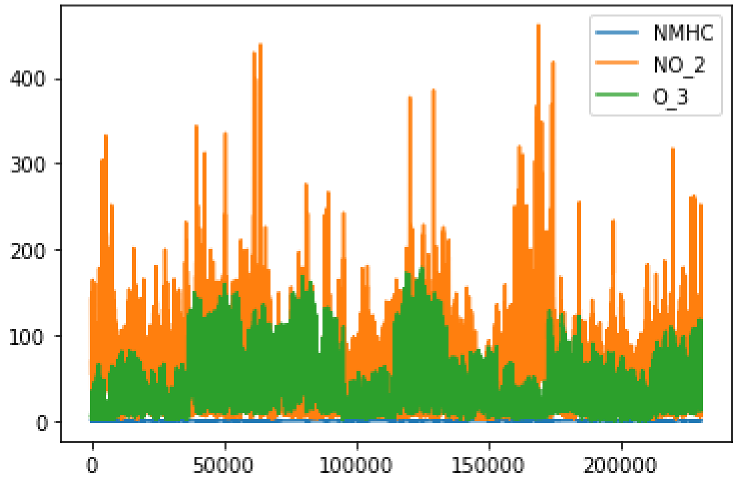

Bar chart

In [11]:

```
b=data[0:50]
```

In [12]:

```
b.plot.bar()
```

Out[12]:

<AxesSubplot:>

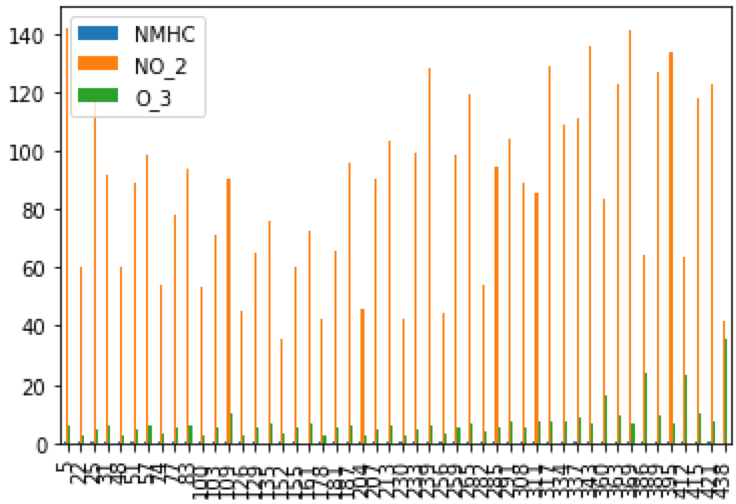

Histogram

In [13]:

```
data.plot.hist()
```

Out[13]:

<AxesSubplot:ylabel='Frequency'>

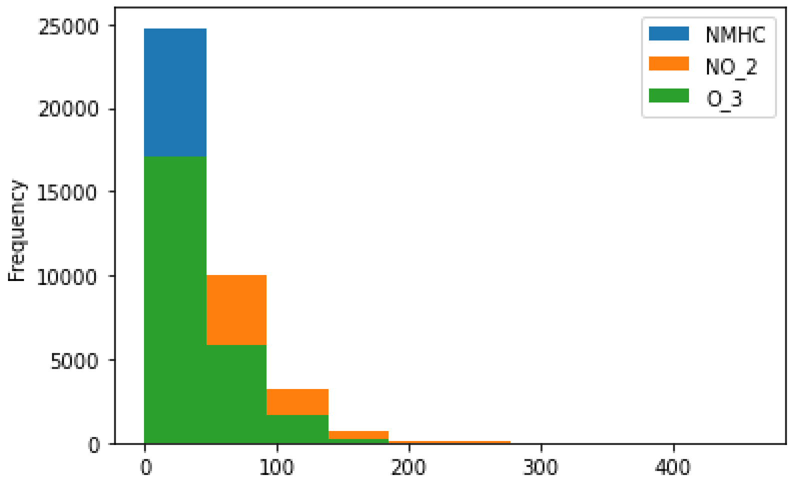

Area chart

In [14]:

```
data.plot.area()
```

Out[14]:

<AxesSubplot:>

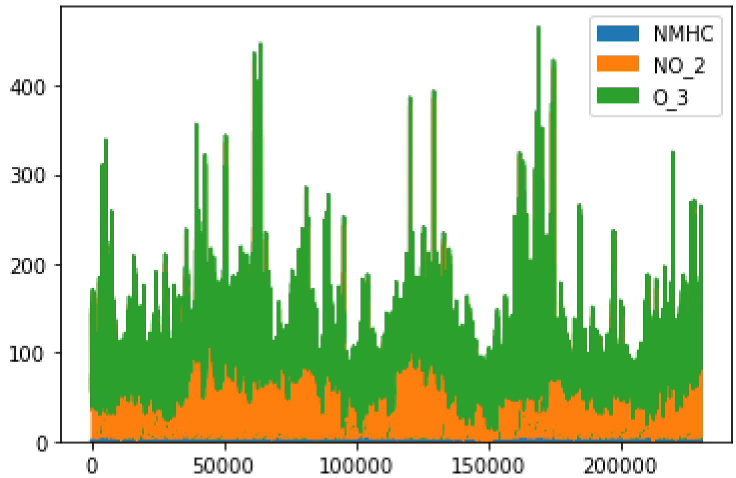

Box chart

In [15]:

```
data.plot.box()
```

Out[15]:

<AxesSubplot:>

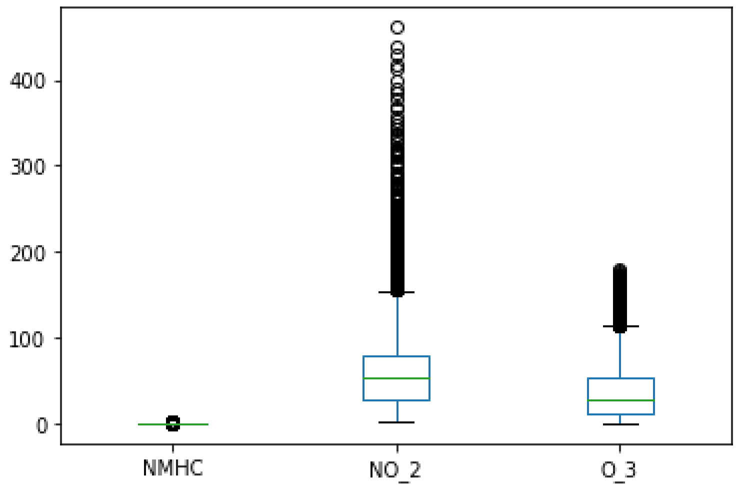

Pie chart

In [17]:

```
b.plot.pie(y='O_3' )
```

Out[17]:

```
<AxesSubplot:ylabel='O_3'>
```

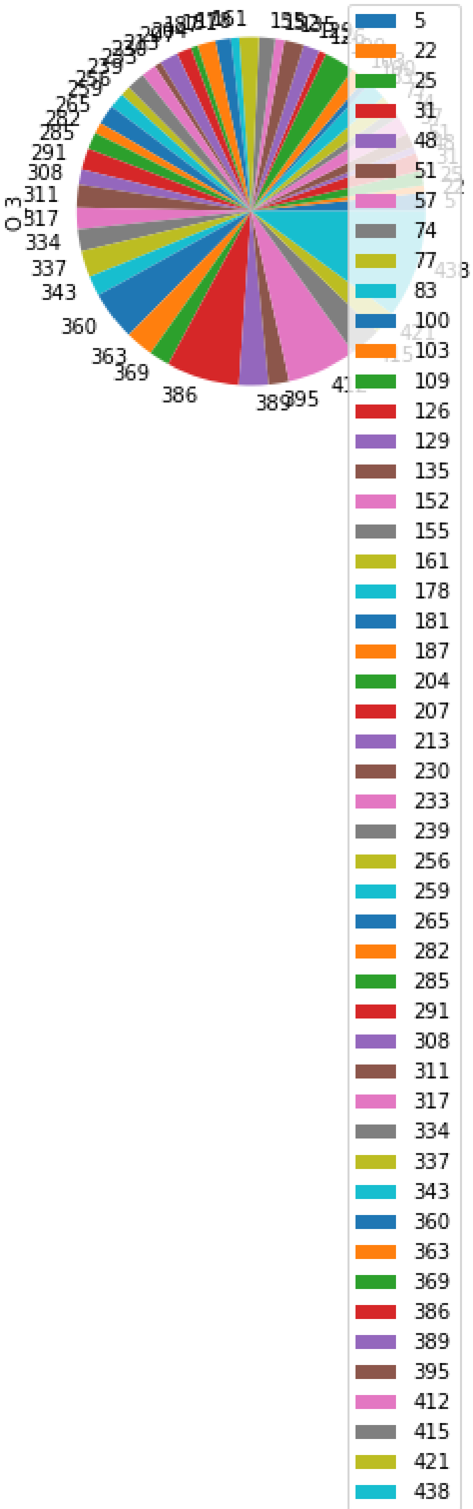

Scatter chart

In [18]:

```
data.plot.scatter(x='NMHC' ,y='O_3')
```

Out[18]:

```
<AxesSubplot:xlabel='NMHC', ylabel='O_3'>
```

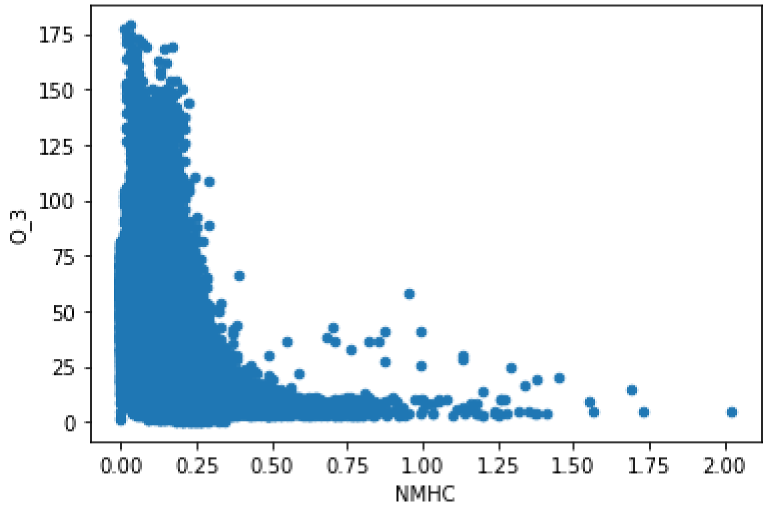

In [19]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 24758 entries, 5 to 230567
Data columns (total 17 columns):
 #   Column   Non-Null Count  Dtype  
---  ------   --------------  -----  
 0   date     24758 non-null  object 
 1   BEN      24758 non-null  float64
 2   CO       24758 non-null  float64
 3   EBE      24758 non-null  float64
 4   MXY      24758 non-null  float64
 5   NMHC     24758 non-null  float64
 6   NO_2     24758 non-null  float64
 7   NOx      24758 non-null  float64
 8   OXY      24758 non-null  float64
 9   O_3      24758 non-null  float64
 10  PM10     24758 non-null  float64
 11  PM25     24758 non-null  float64
 12  PXY      24758 non-null  float64
 13  SO_2     24758 non-null  float64
 14  TCH      24758 non-null  float64
 15  TOL      24758 non-null  float64
 16  station  24758 non-null  int64  
dtypes: float64(15), int64(1), object(1)
memory usage: 3.4+ MB
```

In [20]:

```
df.describe()
```

Out[20]:

	BEN	CO	EBE	MXY	NMHC	NO_2
count	24758.000000	24758.000000	24758.000000	24758.000000	24758.000000	24758.000000
mean	1.350624	0.600713	1.824534	3.835034	0.176546	58.333481
std	1.541636	0.419048	1.868939	4.069036	0.126683	40.529382
min	0.110000	0.000000	0.170000	0.150000	0.000000	1.680000
25%	0.450000	0.360000	0.810000	1.060000	0.100000	28.450001
50%	0.850000	0.500000	1.130000	2.500000	0.150000	52.959999
75%	1.680000	0.720000	2.160000	5.090000	0.220000	79.347498
max	45.430000	7.250000	57.799999	66.900002	2.020000	461.299988

In [21]:

```
df1=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
       'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]
```

EDA AND VISUALIZATION

In [22]:

```
sns.pairplot(df1[0:50])
```

Out[22]:

<seaborn.axisgrid.PairGrid at 0x21a7763ad30>

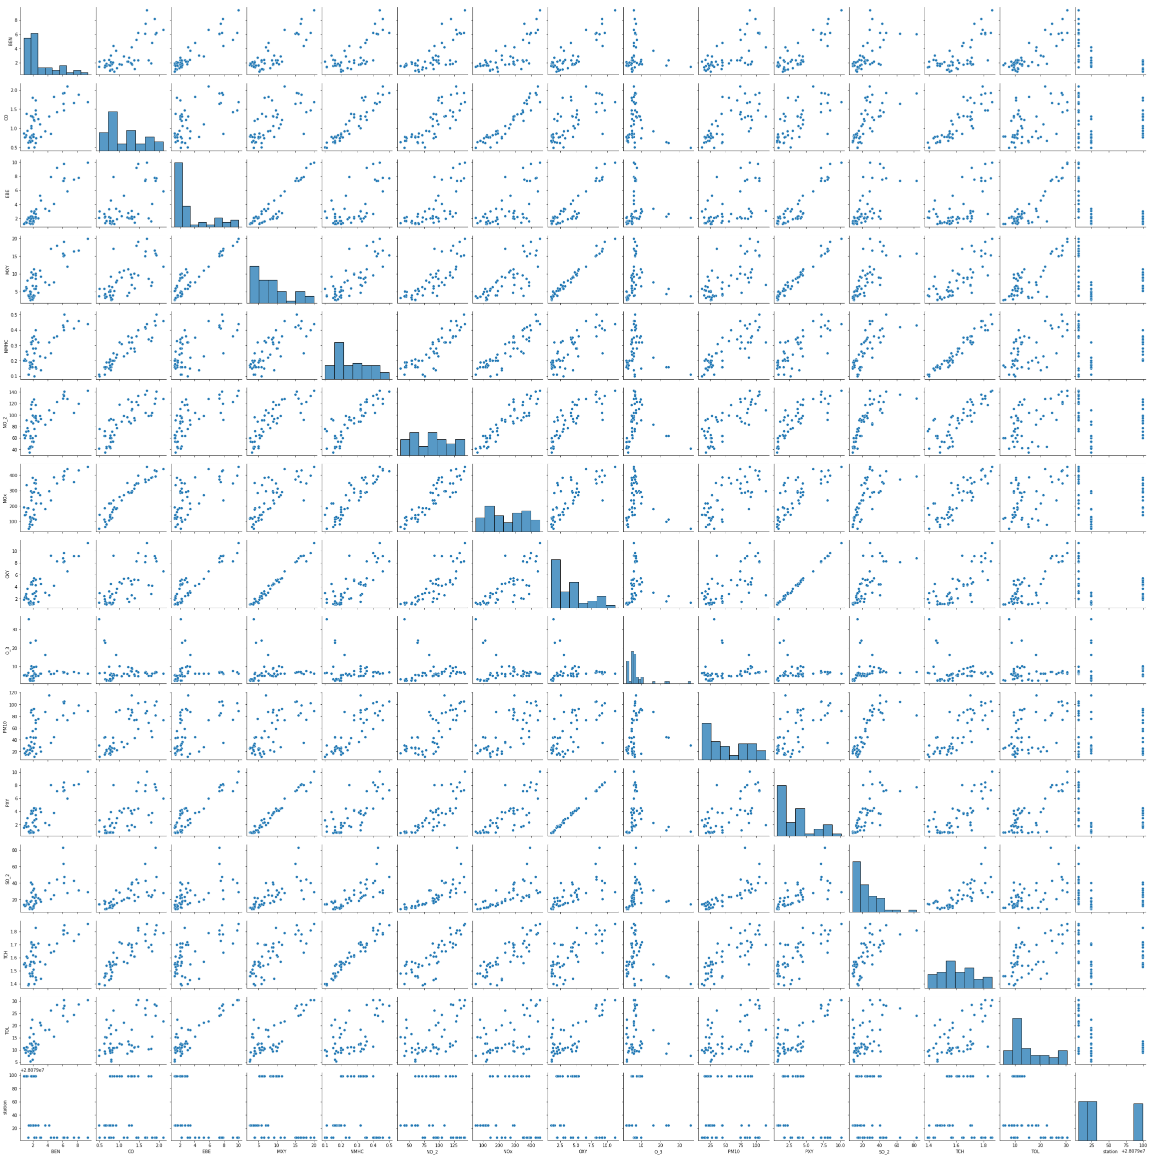

In [23]:

```
sns.distplot(df1['O_3'])
```

```
C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557:
FutureWarning: `distplot` is a deprecated function and will be removed in
a future version. Please adapt your code to use either `displot` (a figure
-level function with similar flexibility) or `histplot` (an axes-level fun
ction for histograms).
  warnings.warn(msg, FutureWarning)
```

Out[23]:

<AxesSubplot:xlabel='O_3', ylabel='Density'>

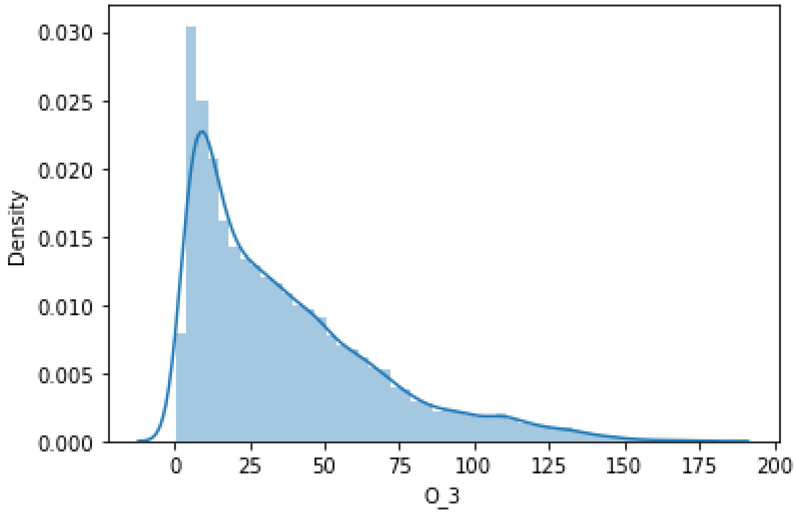

In [24]:

```
sns.heatmap(df1.corr())
```

Out[24]:

<AxesSubplot:>

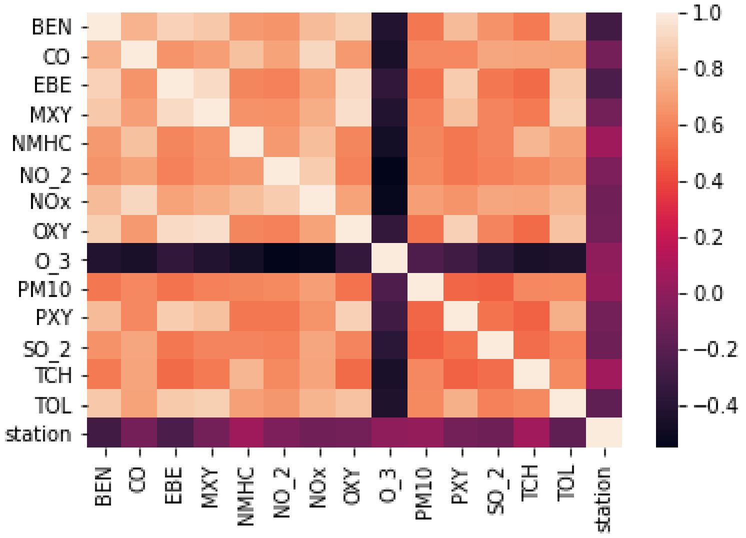

TO TRAIN THE MODEL AND MODEL BULDING

In [67]:

```
x=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY','O_3',
       'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
y=df['station']
```

In [68]:

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

Linear Regression

In [69]:

```
from sklearn.linear_model import LinearRegression
lr=LinearRegression()
lr.fit(x_train,y_train)
```

Out[69]:

```
LinearRegression()
```

In [70]:

```
lr.intercept_
```

Out[70]:

```
28079021.191809315
```

In [71]:

```
coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
coeff
```

Out[71]:

	Co-efficient
BEN	-18.378990
CO	-9.956654
EBE	-23.531143
MXY	4.689670
NMHC	126.649643
NO_2	-0.016238
NOx	-0.004541
OXY	15.492118
O_3	-0.053784
PM10	0.131450
PXY	5.674572
SO_2	-0.652901
TCH	17.650518
TOL	-0.517190

In [72]:

```
prediction =lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[72]:

```
<matplotlib.collections.PathCollection at 0x21a07f864f0>
```

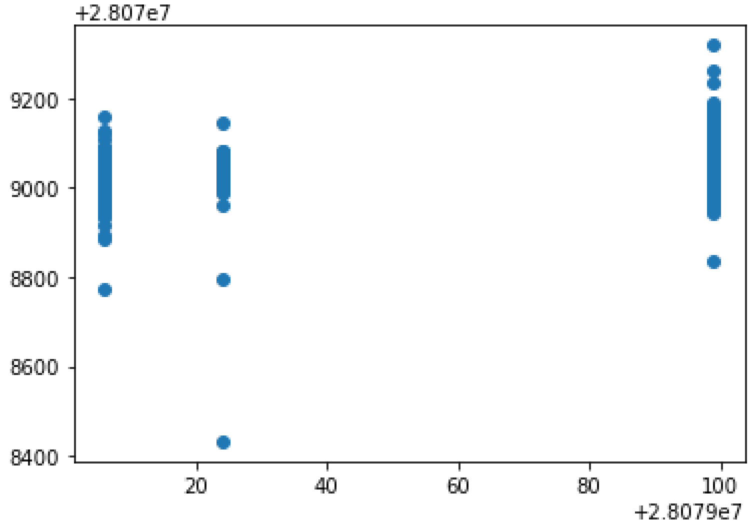

ACCURACY

In [73]:

```
lr.score(x_test,y_test)
```

Out[73]:

```
0.38347413896089644
```

In [74]:

```
lr.score(x_train,y_train)
```

Out[74]:

```
0.3972450971871986
```

Ridge and Lasso

In [75]:

```
from sklearn.linear_model import Ridge,Lasso
```

In [76]:

```
rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

Out[76]:

```
Ridge(alpha=10)
```

Accuracy(Ridge)

In [77]:

```
rr.score(x_test,y_test)
```

Out[77]:

```
0.3820493845006965
```

In [78]:

```
rr.score(x_train,y_train)
```

Out[78]:

```
0.39661106884033615
```

In [79]:

```
la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

Out[79]:

```
Lasso(alpha=10)
```

In [81]:

```
la.score(x_test,y_test)
```

Out[81]:

```
0.05647819697010681
```

Accuracy(Lasso)

In [80]:

```
la.score(x_train,y_train)
```

Out[80]:

```
0.0623362480802043 26
```

Accuracy(Elastic Net)

In [82]:

```
from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

Out[82]:

```
ElasticNet()
```

In [83]:

```
en.coef_
```

Out[83]:

```
array([-8.55559241e+00,  0.00000000e+00, -9.05278481e+00,  3.41965661e+00,
        4.01206179e-01, -3.00597081e-03,  5.10526761e-03,  3.44982093e+00,
       -1.23661221e-01,  2.93481571e-01,  2.43724729e+00, -4.22212620e-01,
        5.28019988e-01, -9.96915988e-01])
```

In [84]:

```
en.intercept_
```

Out[84]:

```
28079052.093522523
```

In [85]:

```
prediction=en.predict(x_test)
```

In [86]:

```
en.score(x_test,y_test)
```

Out[86]:

```
0.2328938395355542
```

Evaluation Metrics

In [87]:

```
from sklearn import metrics
print(metrics.mean_absolute_error(y_test,prediction))
print(metrics.mean_squared_error(y_test,prediction))
print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
32.338665036423414
1266.5452114192653
35.58855450027811
```

Logistic Regression

In [88]:

```
from sklearn.linear_model import LogisticRegression
```

In [89]:

```
feature_matrix=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
       'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
target_vector=df[ 'station']
```

In [90]:

```
feature_matrix.shape
```

Out[90]:

```
(24758, 14)
```

In [91]:

```
target_vector.shape
```

Out[91]:

```
(24758,)
```

In [92]:

```
from sklearn.preprocessing import StandardScaler
```

In [93]:

```
fs=StandardScaler().fit_transform(feature_matrix)
```

In [94]:

```
logr=LogisticRegression(max_iter=10000)
logr.fit(fs,target_vector)
```

Out[94]:

```
LogisticRegression(max_iter=10000)
```

In [95]:

```
observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13,14]]
```

In [96]:

```
prediction=logr.predict(observation)
print(prediction)
```

```
[28079099]
```

In [97]:

```
logr.classes_
```

Out[97]:

```
array([28079006, 28079024, 28079099], dtype=int64)
```

In [98]:

```
logr.score(fs,target_vector)
```

Out[98]:

```
0.8741416915744405
```

In [99]:

```
logr.predict_proba(observation)[0][0]
```

Out[99]:

```
3.5557727473608076e-15
```

In [100]:

```
logr.predict_proba(observation)
```

Out[100]:

```
array([[3.55577275e-15, 7.80743173e-29, 1.00000000e+00]])
```

Random Forest

In [101]:

```
from sklearn.ensemble import RandomForestClassifier
```

In [102]:

```
rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

Out[102]:

```
RandomForestClassifier()
```

In [103]:

```
parameters={'max_depth':[1,2,3,4,5],
            'min_samples_leaf':[5,10,15,20,25],
            'n_estimators':[10,20,30,40,50]
}
```

In [104]:

```
from sklearn.model_selection import GridSearchCV
grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")
grid_search.fit(x_train,y_train)
```

Out[104]:

```
GridSearchCV(cv=2, estimator=RandomForestClassifier(),
             param_grid={'max_depth': [1, 2, 3, 4, 5],
                         'min_samples_leaf': [5, 10, 15, 20, 25],
                         'n_estimators': [10, 20, 30, 40, 50]},
             scoring='accuracy')
```

In [105]:

```
grid_search.best_score_
```

Out[105]:

```
0.8757068667051355
```

In [106]:

```
rfc_best=grid_search.best_estimator_
```

In [107]:

```
from sklearn.tree import plot_tree

plt.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['a','b','c','d'],f
```

```
[113, 64, 159]\nclass = c'),
 Text(4394.25, 181.19999999999982, 'gini = 0.42\nsamples = 850\nvalue =
[1015, 37, 365]\nclass = a')]
```

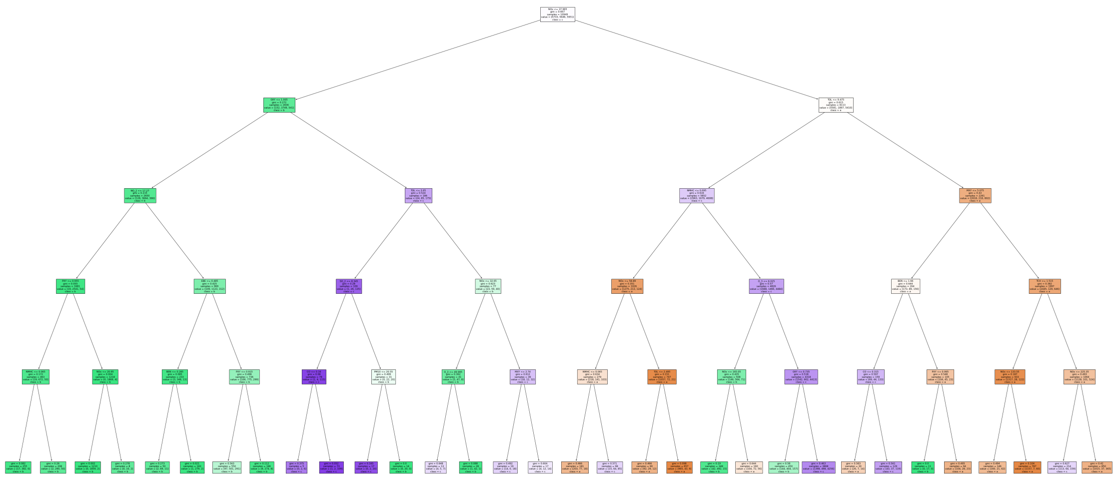

Conclusion

Accuracy

Linear Regression:0.3972450971871986

Ridge Regression:0.39661106884033615

Lasso Regression:0.062336248080204326

ElasticNet Regression:0.2328938395355542

Logistic Regression:0.8741416915744405

Random Forest:0.8757068667051355

Random Forest is suitable for this dataset