

# Importing Libraries

In [1]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

# Importing Datasets

In [2]:

```
df=pd.read_csv(r"C:\Users\user\Downloads\C10_air\csvs_per_year\csvs(Dataset)\madrid_2010.
df
```

Out[2]:

	date	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3	
0	2010-03-01 01:00:00	NaN	0.29	NaN	NaN	NaN	25.090000	29.219999	NaN	68.930000	
1	2010-03-01 01:00:00	NaN	0.27	NaN	NaN	NaN	24.879999	30.040001	NaN	NaN	
2	2010-03-01 01:00:00	NaN	0.28	NaN	NaN	NaN	17.410000	20.540001	NaN	72.120003	
3	2010-03-01 01:00:00	0.38	0.24	1.74	NaN	0.05	15.610000	21.080000	NaN	72.970001	19
4	2010-03-01 01:00:00	0.79	NaN	1.32	NaN	NaN	21.430000	26.070000	NaN	NaN	24
...	...	...	...	...	...	...	...	...	...	...	
209443	2010-08-01 00:00:00	NaN	0.55	NaN	NaN	NaN	125.000000	219.899994	NaN	25.379999	
209444	2010-08-01 00:00:00	NaN	0.27	NaN	NaN	NaN	45.709999	47.410000	NaN	NaN	51
209445	2010-08-01 00:00:00	NaN	NaN	NaN	NaN	0.24	46.560001	49.040001	NaN	46.250000	
209446	2010-08-01 00:00:00	NaN	NaN	NaN	NaN	NaN	46.770000	50.119999	NaN	77.709999	
209447	2010-08-01 00:00:00	0.92	0.43	0.71	NaN	0.25	76.330002	88.190002	NaN	52.259998	47

209448 rows × 17 columns

# Data Cleaning and Data Preprocessing

In [3]:

```
df=df.dropna()
```

In [4]:

```
df.columns
```

Out[4]:

```
Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
      'PM10', 'PM25', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],
      dtype='object')
```

In [5]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 6666 entries, 11 to 191927
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        6666 non-null   object
1   BEN         6666 non-null   float64
2   CO          6666 non-null   float64
3   EBE         6666 non-null   float64
4   MXY         6666 non-null   float64
5   NMHC        6666 non-null   float64
6   NO_2        6666 non-null   float64
7   NOx         6666 non-null   float64
8   OXY         6666 non-null   float64
9   O_3         6666 non-null   float64
10  PM10        6666 non-null   float64
11  PM25        6666 non-null   float64
12  PXY         6666 non-null   float64
13  SO_2        6666 non-null   float64
14  TCH         6666 non-null   float64
15  TOL         6666 non-null   float64
16  station     6666 non-null   int64
dtypes: float64(15), int64(1), object(1)
memory usage: 937.4+ KB
```

In [11]:

```
data=df[['BEN', 'TOL', 'TCH']]
data
```

Out[11]:

	BEN	TOL	TCH
11	0.78	1.99	1.55
23	0.70	2.62	1.48
35	0.58	0.84	1.54
47	0.33	1.21	1.44
59	0.38	0.49	1.54
...	...	...	...
191879	0.60	2.94	1.34
191891	0.41	1.11	1.31
191903	0.57	2.95	1.36
191915	0.34	1.09	1.32
191927	0.43	2.80	1.38

6666 rows × 3 columns

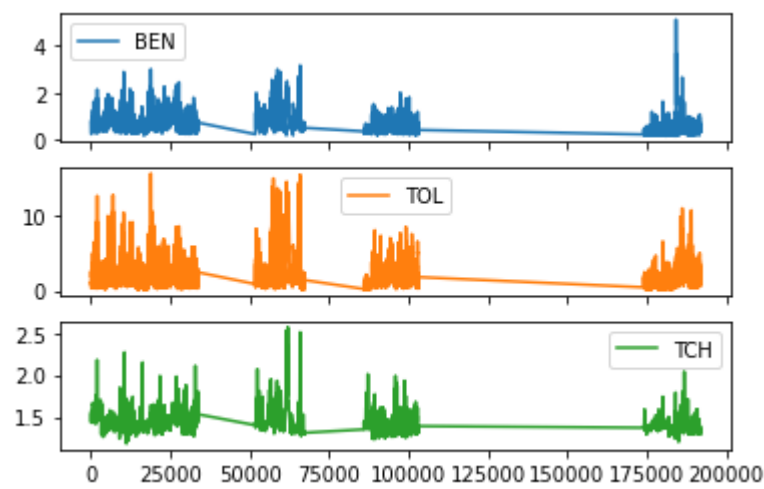
## Line chart

In [12]:

```
data.plot.line(subplots=True)
```

Out[12]:

array([<AxesSubplot:>, <AxesSubplot:>, <AxesSubplot:>], dtype=object)




## Line chart

In [13]:

```
data.plot.line()
```

Out[13]:

&lt;AxesSubplot:&gt;



## Bar chart

In [14]:


```
b=data[0:50]
```

In [15]:

```
b.plot.bar()
```

Out[15]:

&lt;AxesSubplot:&gt;




## Histogram

In [16]:

```
data.plot.hist()
```

Out[16]:

<AxesSubplot:ylabel='Frequency'>




## Area chart

In [17]:

```
data.plot.area()
```

Out[17]:

<AxesSubplot:>



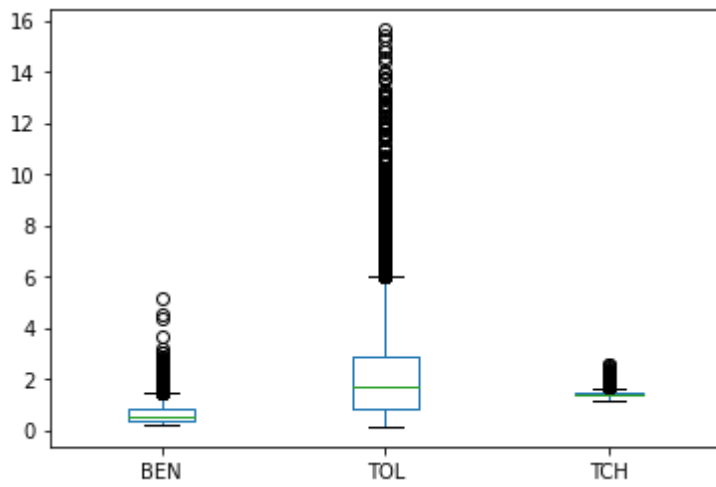
## Box chart

In [18]:

```
data.plot.box()
```

Out[18]:

<AxesSubplot:>




## Pie chart

In [19]:

```
b.plot.pie(y='BEN' )
```

Out[19]:

<AxesSubplot:ylabel='BEN'>



# Scatter chart




In [20]:

```
data.plot.scatter(x='BEN' ,y='TOL')
```

Out[20]:

<AxesSubplot:xlabel='BEN', ylabel='TOL'>



In [21]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 6666 entries, 11 to 191927
Data columns (total 17 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   date        6666 non-null   object  
 1   BEN         6666 non-null   float64 
 2   CO          6666 non-null   float64 
 3   EBE         6666 non-null   float64 
 4   MXY         6666 non-null   float64 
 5   NMHC        6666 non-null   float64 
 6   NO_2        6666 non-null   float64 
 7   NOx         6666 non-null   float64 
 8   OXY         6666 non-null   float64 
 9   O_3         6666 non-null   float64 
10  PM10        6666 non-null   float64 
11  PM25        6666 non-null   float64 
12  PXY         6666 non-null   float64 
13  SO_2        6666 non-null   float64 
14  TCH         6666 non-null   float64 
15  TOL         6666 non-null   float64 
16  station     6666 non-null   int64   
dtypes: float64(15), int64(1), object(1)
memory usage: 937.4+ KB
```

In [22]:

```
df.describe()
```

Out[22]:

	BEN	CO	EBE	MXY	NMHC	NO_2	
count	6666.000000	6666.000000	6666.000000	6666.000000	6666.000000	6666.000000	6666.000000
mean	0.648425	0.296280	0.840585	0.839959	0.243378	33.888744	47.500000
std	0.395346	0.133296	0.508031	0.382263	0.115730	23.465169	41.200000
min	0.170000	0.090000	0.140000	0.110000	0.000000	1.290000	2.700000
25%	0.380000	0.200000	0.470000	0.590000	0.180000	15.752500	19.400000
50%	0.540000	0.260000	0.755000	1.000000	0.220000	29.320000	36.700000
75%	0.810000	0.340000	1.000000	1.000000	0.280000	47.657500	62.100000
max	5.110000	1.590000	5.190000	6.810000	0.930000	133.399994	409.200000

In [23]:

```
df1=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',  
        'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]
```

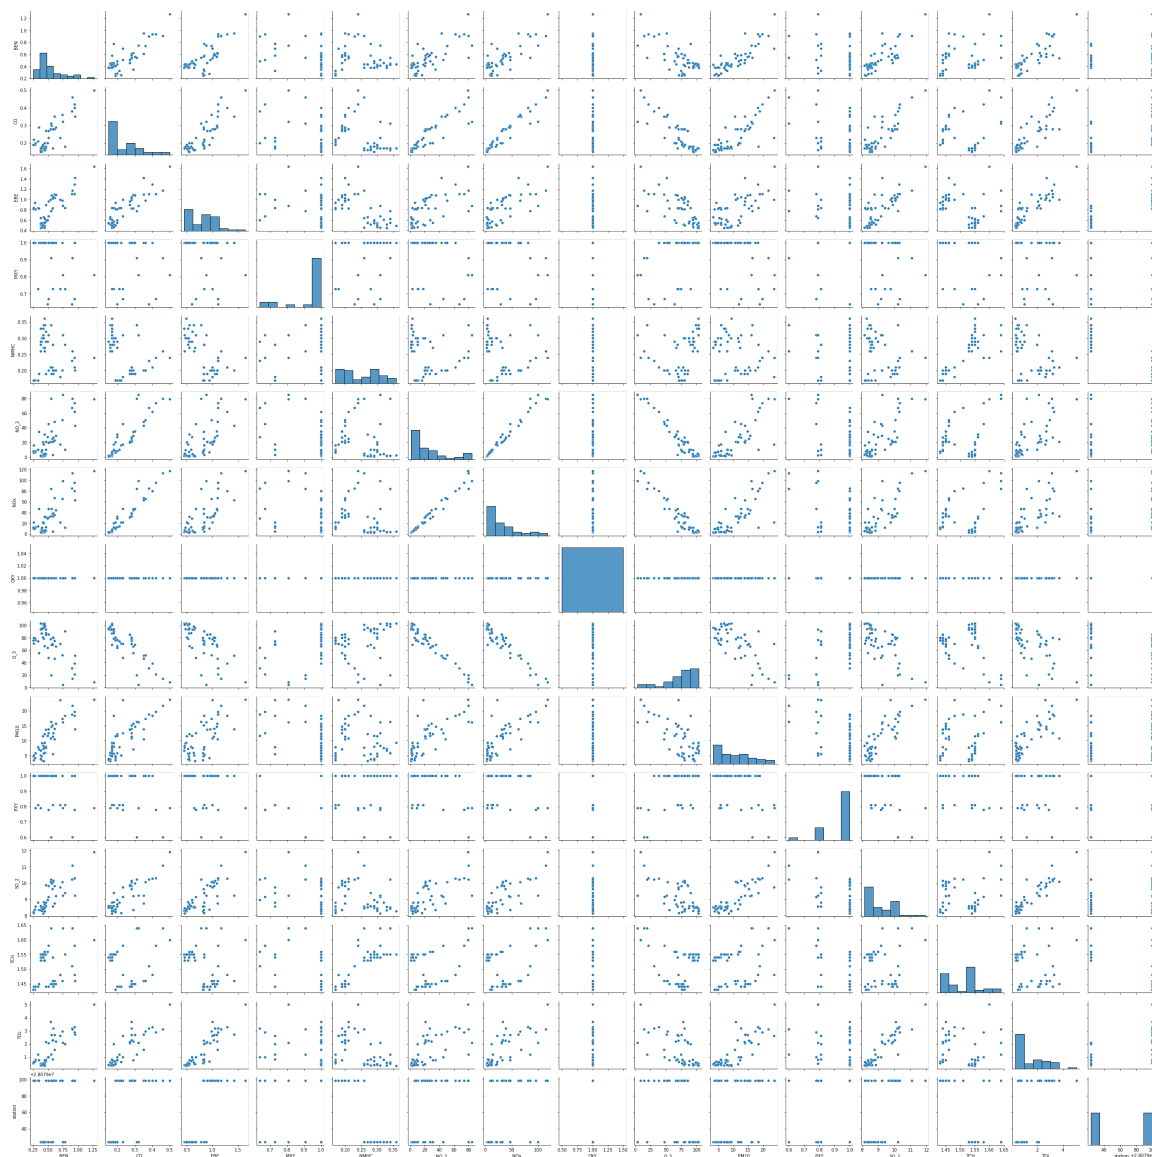
# EDA AND VISUALIZATION

In [24]:

```
sns.pairplot(df1[0:50])
```

Out[24]:

&lt;seaborn.axisgrid.PairGrid at 0x1e6db036850&gt;



In [25]:


```
sns.distplot(df1['BEN'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557:  
FutureWarning: `distplot` is a deprecated function and will be removed in  
a future version. Please adapt your code to use either `displot` (a figure  
-level function with similar flexibility) or `histplot` (an axes-level fun  
ction for histograms).

```
warnings.warn(msg, FutureWarning)
```

Out[25]:

```
<AxesSubplot:xlabel='BEN', ylabel='Density'>
```




In [26]:

```
sns.heatmap(df1.corr())
```

Out[26]:

```
<AxesSubplot:>
```



## TO TRAIN THE MODEL AND MODEL BUILDING

In [111]:

```
x=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',  
      'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]  
y=df['station']
```

In [112]:

```
from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

## Linear Regression

In [113]:

```
from sklearn.linear_model import LinearRegression  
lr=LinearRegression()  
lr.fit(x_train,y_train)
```

Out[113]:

LinearRegression()

In [114]:

```
lr.intercept_
```

Out[114]:

28078940.864442065

In [115]:

```
coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
coeff
```

Out[115]:


	Co-efficient
<b>BEN</b>	-34.991509
<b>CO</b>	169.287030
<b>EBE</b>	16.889120
<b>MXY</b>	-8.950140
<b>NMHC</b>	-75.485838
<b>NO_2</b>	0.312985
<b>NOx</b>	-0.626738
<b>OXY</b>	30.132175
<b>O_3</b>	0.076435
<b>PM10</b>	-0.162293
<b>PXY</b>	-4.783814
<b>SO_2</b>	1.781796
<b>TCH</b>	42.165548
<b>TOL</b>	9.666856

In [116]:

```
prediction =lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[116]:

&lt;matplotlib.collections.PathCollection at 0x1e6ec61a4c0&gt;



## ACCURACY

In [117]:

```
lr.score(x_test,y_test)
```

Out[117]:

0.4179054524604561

In [118]:

```
lr.score(x_train,y_train)
```

Out[118]:

0.4315656379957823

## Ridge and Lasso

In [119]:

```
from sklearn.linear_model import Ridge,Lasso
```

In [120]:

```
rr=Ridge(alpha=10)  
rr.fit(x_train,y_train)
```

Out[120]:

Ridge(alpha=10)

## Accuracy(Ridge)

In [121]:

```
rr.score(x_test,y_test)
```

Out[121]:

0.4056035229335304

In [122]:

```
rr.score(x_train,y_train)
```

Out[122]:

0.4193276233024624

In [123]:

```
la=Lasso(alpha=10)  
la.fit(x_train,y_train)
```

Out[123]:

Lasso(alpha=10)

In [124]:

```
la.score(x_test,y_test)
```

Out[124]:

0.18344759737444682

## Accuracy(Lasso)

In [125]:

```
la.score(x_train,y_train)
```

Out[125]:

0.17829658999839737

## Accuracy(Elastic Net)

In [126]:

```
from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

Out[126]:

ElasticNet()

In [127]:

```
en.coef_
```

Out[127]:

```
array([-0.          ,  0.19545881,  2.94342568, -1.19054764, -1.31923906,
        0.05983388, -0.12082628,  0.40926303, -0.01846025, -0.12404605,
        -0.          ,  2.57206502,  0.          ,  7.04094312])
```

In [128]:

```
en.intercept_
```

Out[128]:

28079025.73235495

In [129]:

```
prediction=en.predict(x_test)
```



In [130]:

```
en.score(x_test,y_test)
```

Out[130]:

0.23954371788037176

## Evaluation Metrics

In [131]:

```
from sklearn import metrics
print(metrics.mean_absolute_error(y_test,prediction))
print(metrics.mean_squared_error(y_test,prediction))
print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

30.817089282844215

1069.284707566054

32.699919075833414

## Logistic Regression

In [132]:

```
from sklearn.linear_model import LogisticRegression
```

In [133]:

```
feature_matrix=df[['BEN', 'CO', 'EBE', 'MXV', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
                  'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
target_vector=df[ 'station']
```

In [134]:

```
feature_matrix.shape
```

Out[134]:

(6666, 14)

In [135]:

```
target_vector.shape
```

Out[135]:

(6666,)

In [136]:

```
from sklearn.preprocessing import StandardScaler
```

In [137]:

```
fs=StandardScaler().fit_transform(feature_matrix)
```

In [138]:

```
logr=LogisticRegression(max_iter=10000)  
logr.fit(fs,target_vector)
```

Out[138]:

```
LogisticRegression(max_iter=10000)
```

In [139]:

```
observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13,14]]
```

In [140]:

```
prediction=logr.predict(observation)  
print(prediction)
```

```
[28079099]
```

In [141]:

```
logr.classes_
```

Out[141]:

```
array([28079024, 28079099], dtype=int64)
```

In [142]:

```
logr.score(fs,target_vector)
```

Out[142]:

```
0.8660366036603661
```

In [143]:

```
logr.predict_proba(observation)[0][0]
```

Out[143]:

```
0.0
```

In [144]:

```
logr.predict_proba(observation)
```

Out[144]:

```
array([[0., 1.]])
```

## Random Forest

In [145]:

```
from sklearn.ensemble import RandomForestClassifier
```

In [146]:

```
rfc=RandomForestClassifier()  
rfc.fit(x_train,y_train)
```

Out[146]:

```
RandomForestClassifier()
```

In [147]:

```
parameters={'max_depth':[1,2,3,4,5],  
            'min_samples_leaf':[5,10,15,20,25],  
            'n_estimators':[10,20,30,40,50]  
}
```

In [148]:

```
from sklearn.model_selection import GridSearchCV  
grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")  
grid_search.fit(x_train,y_train)
```

Out[148]:

```
GridSearchCV(cv=2, estimator=RandomForestClassifier(),  
             param_grid={'max_depth': [1, 2, 3, 4, 5],  
                         'min_samples_leaf': [5, 10, 15, 20, 25],  
                         'n_estimators': [10, 20, 30, 40, 50]},  
             scoring='accuracy')
```

In [149]:

```
grid_search.best_score_
```

Out[149]:

```
0.9303471924560651
```

In [150]:

```
rfc_best=grid_search.best_estimator_
```

In [151]:

```
from sklearn.tree import plot_tree

plt.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['a','b','c','d'],f
```

Out[151]:

```

[Text(2576.6470588235293, 1993.2, 'CO <= 0.225\ngini = 0.5\nsamples = 2962\nvalue = [2301, 2365]\nnclass = b'),
Text(1542.705882352941, 1630.8000000000002, 'NMHC <= 0.255\ngini = 0.195\nsamples = 989\nvalue = [1357, 167]\nnclass = a'),
Text(722.1176470588234, 1268.4, 'BEN <= 0.305\ngini = 0.347\nsamples = 477\nvalue = [553, 159]\nnclass = a'),
Text(262.5882352941176, 906.0, 'NMHC <= 0.115\ngini = 0.488\nsamples = 136\nvalue = [85, 116]\nnclass = b'),
Text(131.2941176470588, 543.5999999999999, 'gini = 0.0\nsamples = 40\nvalue = [55, 0]\nnclass = a'),
Text(393.88235294117646, 543.5999999999999, 'NOx <= 11.0\ngini = 0.327\nsamples = 96\nvalue = [30, 116]\nnclass = b'),
Text(262.5882352941176, 181.19999999999982, 'gini = 0.227\nsamples = 16\nvalue = [20, 3]\nnclass = a'),
Text(525.1764705882352, 181.19999999999982, 'gini = 0.149\nsamples = 80\nvalue = [10, 113]\nnclass = b'),
Text(1181.6470588235293, 906.0, 'TOL <= 0.835\ngini = 0.154\nsamples = 341\nvalue = [468, 43]\nnclass = a'),
Text(919.0588235294117, 543.5999999999999, 'NO_2 <= 11.855\ngini = 0.023\nsamples = 169\nvalue = [259, 3]\nnclass = a'),
Text(787.7647058823529, 181.19999999999982, 'gini = 0.0\nsamples = 120\nvalue = [184, 0]\nnclass = a'),
Text(1050.3529411764705, 181.19999999999982, 'gini = 0.074\nsamples = 49\nvalue = [75, 3]\nnclass = a'),
Text(1444.2352941176468, 543.5999999999999, 'MXV <= 0.795\ngini = 0.27\nsamples = 1\nvalue = [209, 40]\nnclass = a'),
Text(1312.941176470588, 181.19999999999982, 'gini = 0.165\nsamples = 93\nvalue = [120, 12]\nnclass = a'),
Text(1575.5294117647058, 181.19999999999982, 'gini = 0.64\nsamples = 19\nvalue = [89, 28]\nnclass = a'),
Text(2363.2941176470586, 1268.4, 'TOL <= 2.05\ngini = 0.02\nsamples = 512\nvalue = [804, 0]\nnclass = a'),
Text(2232.0, 906.0, 'NO_2 <= 18.56\ngini = 0.008\nsamples = 496\nvalue = [789, 3]\nnclass = a'),
Text(1969.4117647058822, 543.5999999999999, 'NMHC <= 0.275\ngini = 0.003\nsamples = 413\nvalue = [649, 1]\nnclass = a'),
Text(1838.1176470588234, 181.19999999999982, 'gini = 0.05\nsamples = 30\nvalue = [38, 1]\nnclass = a'),
Text(2021.7647058823529, 181.19999999999982, 'gini = 0.0\nsamples = 383\nvalue = [611, 0]\nnclass = a'),
Text(2494.5882352941176, 543.5999999999999, 'PM10 <= 9.79\ngini = 0.028\nsamples = 83\nvalue = [139, 2]\nnclass = a'),
Text(2565.7647058823529, 181.19999999999982, 'gini = 0.102\nsamples = 28\nvalue = [35, 2]\nnclass = a'),
Text(2625.882352941176, 181.19999999999982, 'gini = 0.0\nsamples = 55\nvalue = [104, 0]\nnclass = a'),
Text(2494.5882352941176, 906.0, 'gini = 0.363\nsamples = 16\nvalue = [16, 0]\nnclass = a'),
Text(3610.588235294117, 1630.8000000000002, 'O_3 <= 5.67\ngini = 0.42\nsamples = 1973\nvalue = [944, 2198]\nnclass = b'),
Text(3479.2941176470586, 1268.4, 'gini = 0.0\nsamples = 92\nvalue = [165, 0]\nnclass = a'),
Text(541.882352941176, 1268.4, 'NMHC <= 0.305\ngini = 0.386\nsamples = 1881\nvalue = [779, 2198]\nnclass = b'),
Text(3282.3529411764703, 906.0, 'CO <= 0.265\ngini = 0.296\nsamples = 1675\nvalue = [478, 2174]\nnclass = b'),
Text(3019.7647058823527, 543.5999999999999, 'EBE <= 0.395\ngini = 0.466\nsamples = 598\nvalue = [226, 586]\nnclass = b'),
Text(2888.4705882352937, 181.19999999999982, 'gini = 0.44\nsamples = 70\nvalue = [70, 34]\nnclass = a'),
Text(3151.0588235294117, 181.19999999999982, 'gini = 0.426\nsamples = 328\nvalue = [158, 321]\nnclass = b')
]

```

**Conclusion**

**Accuracy**

**Linear Regression: 0.4315656379957823**

**Ridge Regression: 0.4193276233024624**

**Lasso Regression: 0.1782965899839737**

**ElasticNet Regression: 0.23954371788037476**

**Logistic Regression: 0.8660366036603661**

**Random Forest: 0.9303471924560651**

```
\nvalue = [156, 352]\nnclass = b'),  
Text(3544.941176470588, 543.5999999999999, 'TOL <= 1.535\ngini = 0.217\nsamples = 1277\nvalue = [252, 1788]\nnclass = b'),  
Text(3413.6470588235293, 181.19999999999982, 'gini = 0.476\nsamples = 150\nvalue = [97, 152]\nnclass = b'),  
Text(3676.235294117647, 181.19999999999982, 'gini = 0.158\nsamples = 1127\nvalue = [155, 1636]\nnclass = b'),  
Text(4201.411764705882, 906.0, 'EBE <= 1.52\ngini = 0.137\nsamples = 206\nvalue = [301, 24]\nnclass = a'),  
Text(4070.117647058823, 543.5999999999999, 'NO_2 <= 24.94\ngini = 0.02\nsamples = 100\nvalue = [304, 24]\nnclass = a')
```

**Random Forest is suitable for this dataset**