

Importing Libraries

In [1]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

Importing Datasets

In [2]:

```
df=pd.read_csv(r"C:\Users\user\Downloads\C10_air\csvs_per_year\csvs(Dataset)\madrid_2011.
df
```

Out[2]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOI
0	2011-11-01 01:00:00	NaN	1.0	NaN	NaN	154.0	84.0	NaN	NaN	NaN	6.0	NaN	NaN
1	2011-11-01 01:00:00	2.5	0.4	3.5	0.26	68.0	92.0	3.0	40.0	24.0	9.0	1.54	8.0
2	2011-11-01 01:00:00	2.9	NaN	3.8	NaN	96.0	99.0	NaN	NaN	NaN	NaN	NaN	7.0
3	2011-11-01 01:00:00	NaN	0.6	NaN	NaN	60.0	83.0	2.0	NaN	NaN	NaN	NaN	NaN
4	2011-11-01 01:00:00	NaN	NaN	NaN	NaN	44.0	62.0	3.0	NaN	NaN	3.0	NaN	NaN
...
209923	2011-09-01 00:00:00	NaN	0.2	NaN	NaN	5.0	19.0	44.0	NaN	NaN	NaN	NaN	NaN
209924	2011-09-01 00:00:00	NaN	0.1	NaN	NaN	6.0	29.0	NaN	11.0	NaN	7.0	NaN	NaN
209925	2011-09-01 00:00:00	NaN	NaN	NaN	0.23	1.0	21.0	28.0	NaN	NaN	NaN	1.44	NaN
209926	2011-09-01 00:00:00	NaN	NaN	NaN	NaN	3.0	15.0	48.0	NaN	NaN	NaN	NaN	NaN
209927	2011-09-01 00:00:00	NaN	NaN	NaN	NaN	4.0	33.0	38.0	13.0	NaN	NaN	NaN	NaN

209928 rows × 14 columns



Data Cleaning and Data Preprocessing

In [3]:

```
df=df.dropna()
```

In [4]:

```
df.columns
```

Out[4]:

```
Index(['date', 'BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'P  
M25',  
      'SO_2', 'TCH', 'TOL', 'station'],  
      dtype='object')
```

In [5]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 16460 entries, 1 to 209910  
Data columns (total 14 columns):  
#   Column      Non-Null Count  Dtype  
---  ---  
0   date        16460 non-null  object  
1   BEN         16460 non-null  float64  
2   CO          16460 non-null  float64  
3   EBE         16460 non-null  float64  
4   NMHC        16460 non-null  float64  
5   NO          16460 non-null  float64  
6   NO_2        16460 non-null  float64  
7   O_3         16460 non-null  float64  
8   PM10        16460 non-null  float64  
9   PM25        16460 non-null  float64  
10  SO_2        16460 non-null  float64  
11  TCH         16460 non-null  float64  
12  TOL         16460 non-null  float64  
13  station     16460 non-null  int64  
dtypes: float64(12), int64(1), object(1)  
memory usage: 1.9+ MB
```

In [6]:

```
data=df[['BEN', 'TOL', 'TCH']]
data
```

Out[6]:

	BEN	TOL	TCH
1	2.5	8.7	1.54
6	0.7	1.7	1.36
25	1.8	7.4	1.71
30	1.0	2.9	1.40
49	1.3	6.2	1.75
...
209862	0.4	0.7	1.26
209881	0.9	4.9	1.34
209886	0.6	0.9	1.26
209905	0.6	3.8	1.32
209910	0.7	0.9	1.25

16460 rows × 3 columns


Line chart

In [7]:

```
data.plot.line(subplots=True)
```

Out[7]:

array([<AxesSubplot:>, <AxesSubplot:>, <AxesSubplot:>], dtype=object)




Line chart

In [8]:

```
data.plot.line()
```

Out[8]:

<AxesSubplot:>



Bar chart

In [9]:


```
b=data[0:50]
```

In [10]:

```
b.plot.bar()
```

Out[10]:

<AxesSubplot:>




Histogram

In [11]:

```
data.plot.hist()
```

Out[11]:

<AxesSubplot:ylabel='Frequency'>




Area chart

In [12]:

```
data.plot.area()
```

Out[12]:

<AxesSubplot:>




Box chart

In [13]:

```
data.plot.box()
```

Out[13]:

<AxesSubplot:>




Pie chart

In [14]:

```
b.plot.pie(y='BEN' )
```

Out[14]:

<AxesSubplot:ylabel='BEN'>




Scatter chart

In [15]:

```
data.plot.scatter(x='BEN' ,y='TOL')
```

Out[15]:

<AxesSubplot:xlabel='BEN', ylabel='TOL'>



In [16]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 16460 entries, 1 to 209910
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   date        16460 non-null  object  
 1   BEN         16460 non-null  float64 
 2   CO          16460 non-null  float64 
 3   EBE         16460 non-null  float64 
 4   NMHC        16460 non-null  float64 
 5   NO          16460 non-null  float64 
 6   NO_2        16460 non-null  float64 
 7   O_3         16460 non-null  float64 
 8   PM10        16460 non-null  float64 
 9   PM25        16460 non-null  float64 
10   SO_2        16460 non-null  float64 
11   TCH         16460 non-null  float64 
12   TOL         16460 non-null  float64 
13   station     16460 non-null  int64   
dtypes: float64(12), int64(1), object(1)
memory usage: 1.9+ MB
```

In [17]:

```
df.describe()
```

Out[17]:

	BEN	CO	EBE	NMHC	NO	NO_2
count	16460.000000	16460.000000	16460.000000	16460.000000	16460.000000	16460.000000
mean	0.900680	0.277758	1.471871	0.167043	23.671810	44.583961
std	0.768892	0.206143	1.051004	0.075068	44.362859	31.569185
min	0.100000	0.100000	0.200000	0.010000	1.000000	1.000000
25%	0.500000	0.200000	0.800000	0.120000	2.000000	19.000000
50%	0.700000	0.200000	1.200000	0.160000	7.000000	40.000000
75%	1.100000	0.300000	1.700000	0.200000	25.000000	63.000000
max	9.500000	3.200000	12.800000	0.840000	615.000000	289.000000

In [18]:

```
df1=df[['date', 'BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',  
        'SO_2', 'TCH', 'TOL', 'station']]
```

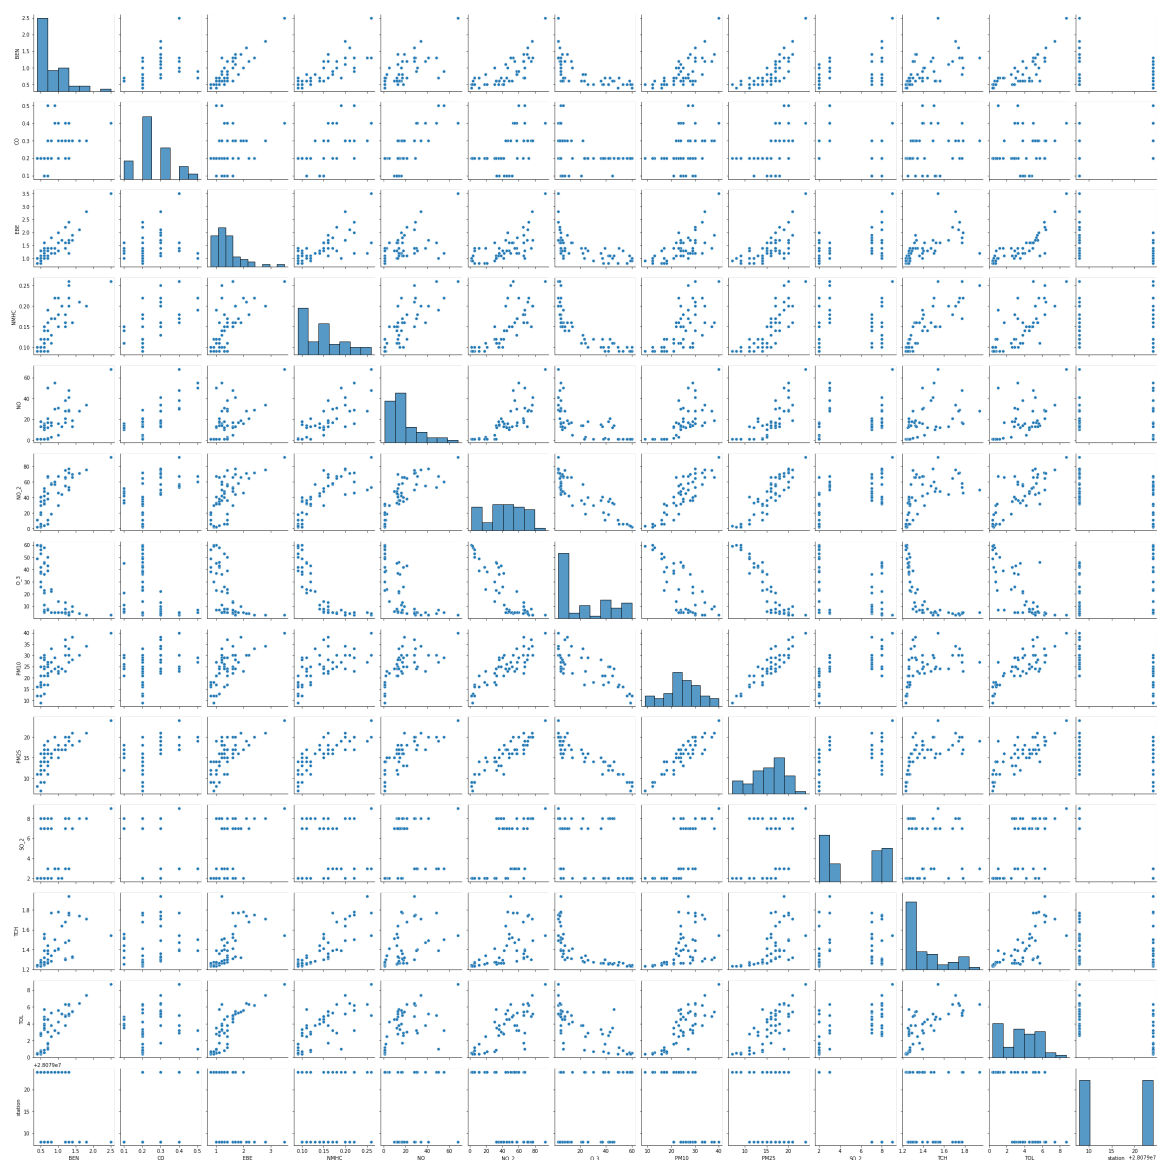
EDA AND VISUALIZATION

In [19]:

```
sns.pairplot(df1[0:50])
```

Out[19]:

<seaborn.axisgrid.PairGrid at 0x1bb2a70d7f0>



In [20]:


```
sns.distplot(df1['BEN'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```

Out[20]:

```
<AxesSubplot:xlabel='BEN', ylabel='Density'>
```




In [21]:

```
sns.heatmap(df1.corr())
```

Out[21]:

```
<AxesSubplot:>
```



TO TRAIN THE MODEL AND MODEL BUILDING

In [22]:

```
x=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',  
      'SO_2', 'TCH', 'TOL']]  
y=df['station']
```

In [23]:

```
from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

Linear Regression

In [24]:

```
from sklearn.linear_model import LinearRegression  
lr=LinearRegression()  
lr.fit(x_train,y_train)
```

Out[24]:

LinearRegression()

In [25]:

```
lr.intercept_
```

Out[25]:

28079015.223056305

In [26]:

```
coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
coeff
```

Out[26]:


	Co-efficient
BEN	3.805170
CO	37.603729
EBE	-1.858668
NMHC	-93.094409
NO	-0.036306
NO_2	-0.089968
O_3	-0.015230
PM10	0.014540
PM25	-0.035677
SO_2	-0.474032
TCH	11.074567
TOL	-0.415248

In [27]:

```
prediction = lr.predict(x_test)
plt.scatter(y_test, prediction)
```

Out[27]:

<matplotlib.collections.PathCollection at 0x1bb36eccac0>



ACCURACY

In [28]:

```
lr.score(x_test, y_test)
```

Out[28]:

0.6176242622614179

In [29]:

```
lr.score(x_train, y_train)
```

Out[29]:

0.6312376538837546

Ridge and Lasso

In [30]:

```
from sklearn.linear_model import Ridge, Lasso
```

In [31]:

```
rr=Ridge(alpha=10)
rr.fit(x_train, y_train)
```

Out[31]:

Ridge(alpha=10)

Accuracy(Ridge)

In [32]:

```
rr.score(x_test,y_test)
```

Out[32]:

```
0.580321441574621
```

In [33]:

```
rr.score(x_train,y_train)
```

Out[33]:

```
0.5979243288685172
```

In [34]:

```
la=Lasso(alpha=10)  
la.fit(x_train,y_train)
```

Out[34]:

```
Lasso(alpha=10)
```

In [35]:

```
la.score(x_test,y_test)
```

Out[35]:

```
0.24603868232736525
```

Accuracy(Lasso)

In [36]:

```
la.score(x_train,y_train)
```

Out[36]:

```
0.23241032291592945
```

Accuracy(Elastic Net)

In [37]:

```
from sklearn.linear_model import ElasticNet  
en=ElasticNet()  
en.fit(x_train,y_train)
```

Out[37]:

```
ElasticNet()
```

In [38]:

```
en.coef_
```

Out[38]:

```
array([ 0.30437437,  0.          , -0.          , -0.          ,  0.05117634,  
       -0.13281545, -0.04260736,  0.02915675,  0.09360197, -0.1808079 ,  
        0.          , -1.00493552])
```

In [39]:

```
en.intercept_
```

Out[39]:

```
28079025.017157324
```

In [40]:

```
prediction=en.predict(x_test)
```

In [41]:

```
en.score(x_test,y_test)
```

Out[41]:

```
0.3502416888161278
```

Evaluation Metrics

In [42]:

```
from sklearn import metrics  
print(metrics.mean_absolute_error(y_test,prediction))  
print(metrics.mean_squared_error(y_test,prediction))  
print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
5.604138864118284
```

```
41.58278557254639
```

```
6.448471568716605
```

Logistic Regression

In [43]:

```
from sklearn.linear_model import LogisticRegression
```

In [44]:

```
feature_matrix=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2', 'O_3',  
                  'PM10', 'SO_2', 'TCH', 'TOL']]  
target_vector=df[ 'station']
```


In [45]:

```
feature_matrix.shape
```

Out[45]:

```
(16460, 10)
```

In [46]:

```
target_vector.shape
```

Out[46]:

```
(16460,)
```

In [47]:

```
from sklearn.preprocessing import StandardScaler
```

In [48]:

```
fs=StandardScaler().fit_transform(feature_matrix)
```

In [49]:

```
logr=LogisticRegression(max_iter=10000)  
logr.fit(fs,target_vector)
```

Out[49]:

```
LogisticRegression(max_iter=10000)
```

In [50]:

```
observation=[[1,2,3,4,5,6,7,8,9,10]]
```

In [51]:

```
prediction=logr.predict(observation)  
print(prediction)
```

```
[28079008]
```

In [52]:

```
logr.classes_
```

Out[52]:

```
array([28079008, 28079024], dtype=int64)
```

In [53]:

```
logr.score(fs,target_vector)
```

Out[53]:

```
0.9237545565006076
```

In [54]:

```
logr.predict_proba(observation)[0][0]
```

Out[54]:

```
0.9999999999999996
```

In [55]:

```
logr.predict_proba(observation)
```

Out[55]:

```
array([[1.00000000e+00, 3.47334507e-15]])
```

Random Forest

In [56]:

```
from sklearn.ensemble import RandomForestClassifier
```

In [57]:

```
rfc=RandomForestClassifier()  
rfc.fit(x_train,y_train)
```

Out[57]:

```
RandomForestClassifier()
```

In [58]:

```
parameters={'max_depth':[1,2,3,4,5],  
            'min_samples_leaf':[5,10,15,20,25],  
            'n_estimators':[10,20,30,40,50]  
}
```

In [59]:

```
from sklearn.model_selection import GridSearchCV  
grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")  
grid_search.fit(x_train,y_train)
```

Out[59]:

```
GridSearchCV(cv=2, estimator=RandomForestClassifier(),  
             param_grid={'max_depth': [1, 2, 3, 4, 5],  
                         'min_samples_leaf': [5, 10, 15, 20, 25],  
                         'n_estimators': [10, 20, 30, 40, 50]},  
             scoring='accuracy')
```

In [60]:

```
grid_search.best_score_
```

Out[60]:

```
0.9395938205172714
```

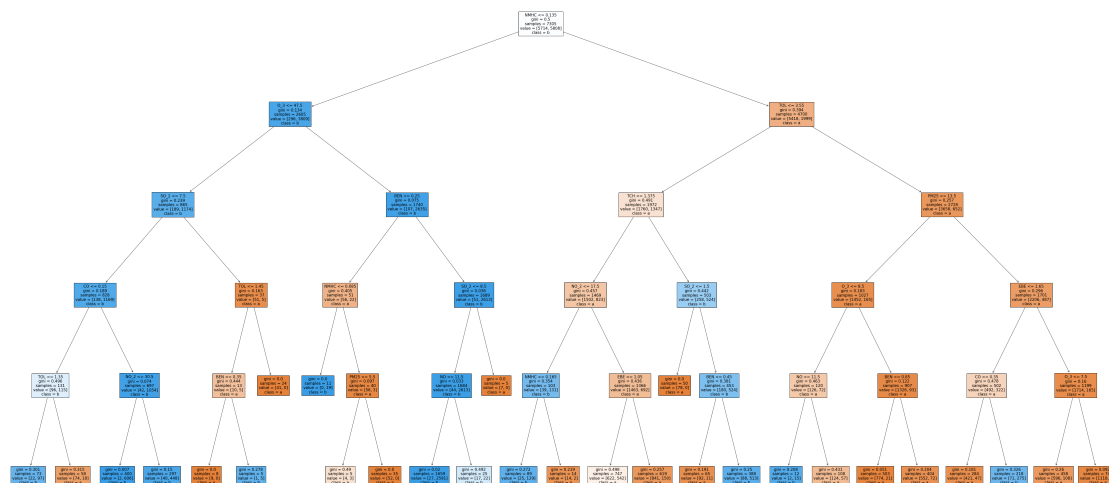
In [61]:

```
rfc_best=grid_search.best_estimator_
```

In [62]:

```
from sklearn.tree import plot_tree

plt.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['a','b','c','d'],f
[596, 108]\nclass = a'),
Text(4374.72, 181.19999999999982, 'gini = 0.092\nsamples = 741\nvalue =
[1118, 57]\nclass = a')]
```



Conclusion

Accuracy

Linear Regression:0.6312376538837546

Ridge Regression:0.5979243288685172

Lasso Regression:0.23241032291592945

ElasticNet Regression:0.3502416888161278

Logistic Regression:0.9999999999999966

Random Forest:0.9395938205172714

Logistic Regression is suitable for this dataset

