

## Importing Libraries

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

## Importing Datasets

```
In [2]: df=pd.read_csv(r"C:\Users\user\Downloads\C10_air\csvs_per_year\csvs(Dataset)\station:
df
```

Out[2]:

	id	name	address	lon	lat	elevation
0	28079004	Pza. de España	Plaza de España	-3.712247	40.423853	635
1	28079008	Escuelas Aguirre	Entre C/ Alcalá y C/ O' Donell	-3.682319	40.421564	670
2	28079011	Avda. Ramón y Cajal	Avda. Ramón y Cajal esq. C/ Príncipe de Vergara	-3.677356	40.451475	708
3	28079016	Arturo Soria	C/ Arturo Soria esq. C/ Vizconde de los Asilos	-3.639233	40.440047	693
4	28079017	Villaverde	C/. Juan Peñalver	-3.713322	40.347139	604
5	28079018	Farolillo	Calle Farolillo - C/Ervigio	-3.731853	40.394781	630
6	28079024	Casa de Campo	Casa de Campo (Terminal del Teleférico)	-3.747347	40.419356	642
7	28079027	Barajas Pueblo	C/. Júpiter, 21 (Barajas)	-3.580031	40.476928	621
8	28079035	Pza. del Carmen	Plaza del Carmen esq. Tres Cruces.	-3.703172	40.419208	659
9	28079036	Moratalaz	Avd. Moratalaz esq. Camino de los Vinateros	-3.645306	40.407947	685
10	28079038	Cuatro Caminos	Avda. Pablo Iglesias esq. C/ Marqués de Lema	-3.707128	40.445544	698
11	28079039	Barrio del Pilar	Avd. Betanzos esq. C/ Monforte de Lemos	-3.711542	40.478228	674
12	28079040	Vallecas	C/ Arroyo del Olivar esq. C/ Río Grande.	-3.651522	40.388153	677
13	28079047	Mendez Alvaro	C/ Juan de Mariana / Pza. Amanecer Mendez Alvaro	-3.686825	40.398114	599
14	28079048	Castellana	C/ Jose Gutierrez Abascal	-3.690367	40.439897	676
15	28079049	Parque del Retiro	Paseo Venezuela- Casa de Vacas	-3.682583	40.414444	662
16	28079050	Plaza Castilla	Plaza Castilla (Canal)	-3.688769	40.465572	728
17	28079054	Ensanche de Vallecas	Avda La Gavia / Avda. Las Suertes	-3.612117	40.372933	627
18	28079055	Urb. Embajada	C/ Riaño (Barajas)	-3.580747	40.462531	618
19	28079056	Pza. Fernández Ladreda	Pza. Fernández Ladreda - Avda. Oporto	-3.718728	40.384964	604
20	28079057	Sanchinarro	C/ Princesa de Eboli esq C/ Maria Tudor	-3.660503	40.494208	700
21	28079058	El Pardo	Avda. La Guardia	-3.774611	40.518058	615
22	28079059	Juan Carlos I	Parque Juan Carlos I (frente oficinas mantenim...	-3.609072	40.465250	660
23	28079060	Tres Olivos	Plaza Tres Olivos	-3.689761	40.500589	715

## Data Cleaning and Data Preprocessing

```
In [3]: df=df.dropna()
```

```
In [4]: df.columns
```

```
Out[4]: Index(['id', 'name', 'address', 'lon', 'lat', 'elevation'], dtype='object')
```

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 24 entries, 0 to 23
Data columns (total 6 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   id          24 non-null    int64
 1   name        24 non-null    object
 2   address     24 non-null    object
 3   lon         24 non-null    float64
 4   lat         24 non-null    float64
 5   elevation   24 non-null    int64
dtypes: float64(2), int64(2), object(2)
memory usage: 1.3+ KB
```

```
In [6]: data=df[['id', 'lon']]
data
```

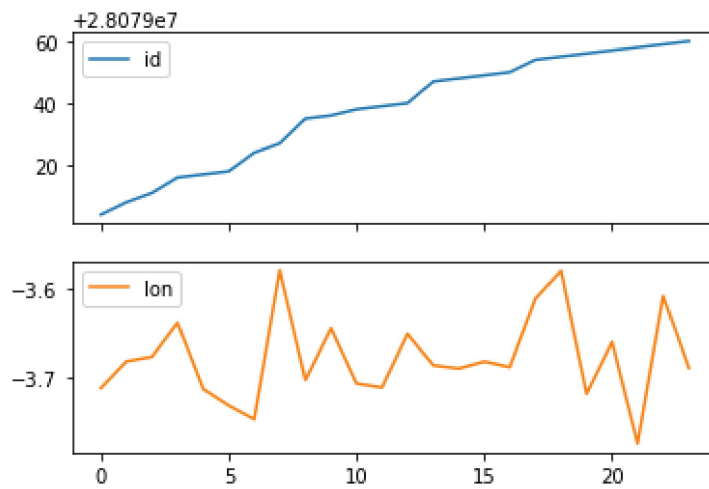
Out[6]:

	id	lon
0	28079004	-3.712247
1	28079008	-3.682319
2	28079011	-3.677356
3	28079016	-3.639233
4	28079017	-3.713322
5	28079018	-3.731853
6	28079024	-3.747347
7	28079027	-3.580031
8	28079035	-3.703172
9	28079036	-3.645306
10	28079038	-3.707128
11	28079039	-3.711542
12	28079040	-3.651522
13	28079047	-3.686825
14	28079048	-3.690367
15	28079049	-3.682583
16	28079050	-3.688769
17	28079054	-3.612117
18	28079055	-3.580747
19	28079056	-3.718728
20	28079057	-3.660503
21	28079058	-3.774611
22	28079059	-3.609072
23	28079060	-3.689761

## Line chart

```
In [7]: data.plot.line(subplots=True)
```


```
Out[7]: array([<AxesSubplot:>, <AxesSubplot:>], dtype=object)
```



## Line chart

```
In [8]: data.plot.line()
```

```
Out[8]: <AxesSubplot:>
```




## Bar chart

```
In [9]: b=data[0:50]
```

```
In [10]: b.plot.bar()
```


```
Out[10]: <AxesSubplot:>
```



## Histogram

```
In [11]: data.plot.hist()
```


```
Out[11]: <AxesSubplot:ylabel='Frequency'>
```



## Area chart

```
In [12]: data.plot.area()
```


```
Out[12]: <AxesSubplot:>
```



## Box chart

```
In [13]: data.plot.box()
```


```
Out[13]: <AxesSubplot:>
```



## Pie chart

```
In [16]: b.plot.pie(y='id' )
```


```
Out[16]: <AxesSubplot:ylabel='id'>
```



## Scatter chart

```
In [17]: data.plot.scatter(x='id' ,y='lon')
```

```
Out[17]: <AxesSubplot:xlabel='id', ylabel='lon'>
```





In [18]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 24 entries, 0 to 23
Data columns (total 6 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   id          24 non-null    int64
 1   name        24 non-null    object
 2   address     24 non-null    object
 3   lon         24 non-null    float64
 4   lat         24 non-null    float64
 5   elevation   24 non-null    int64
dtypes: float64(2), int64(2), object(2)
memory usage: 1.3+ KB
```

In [19]: `df.describe()`

Out[19]:

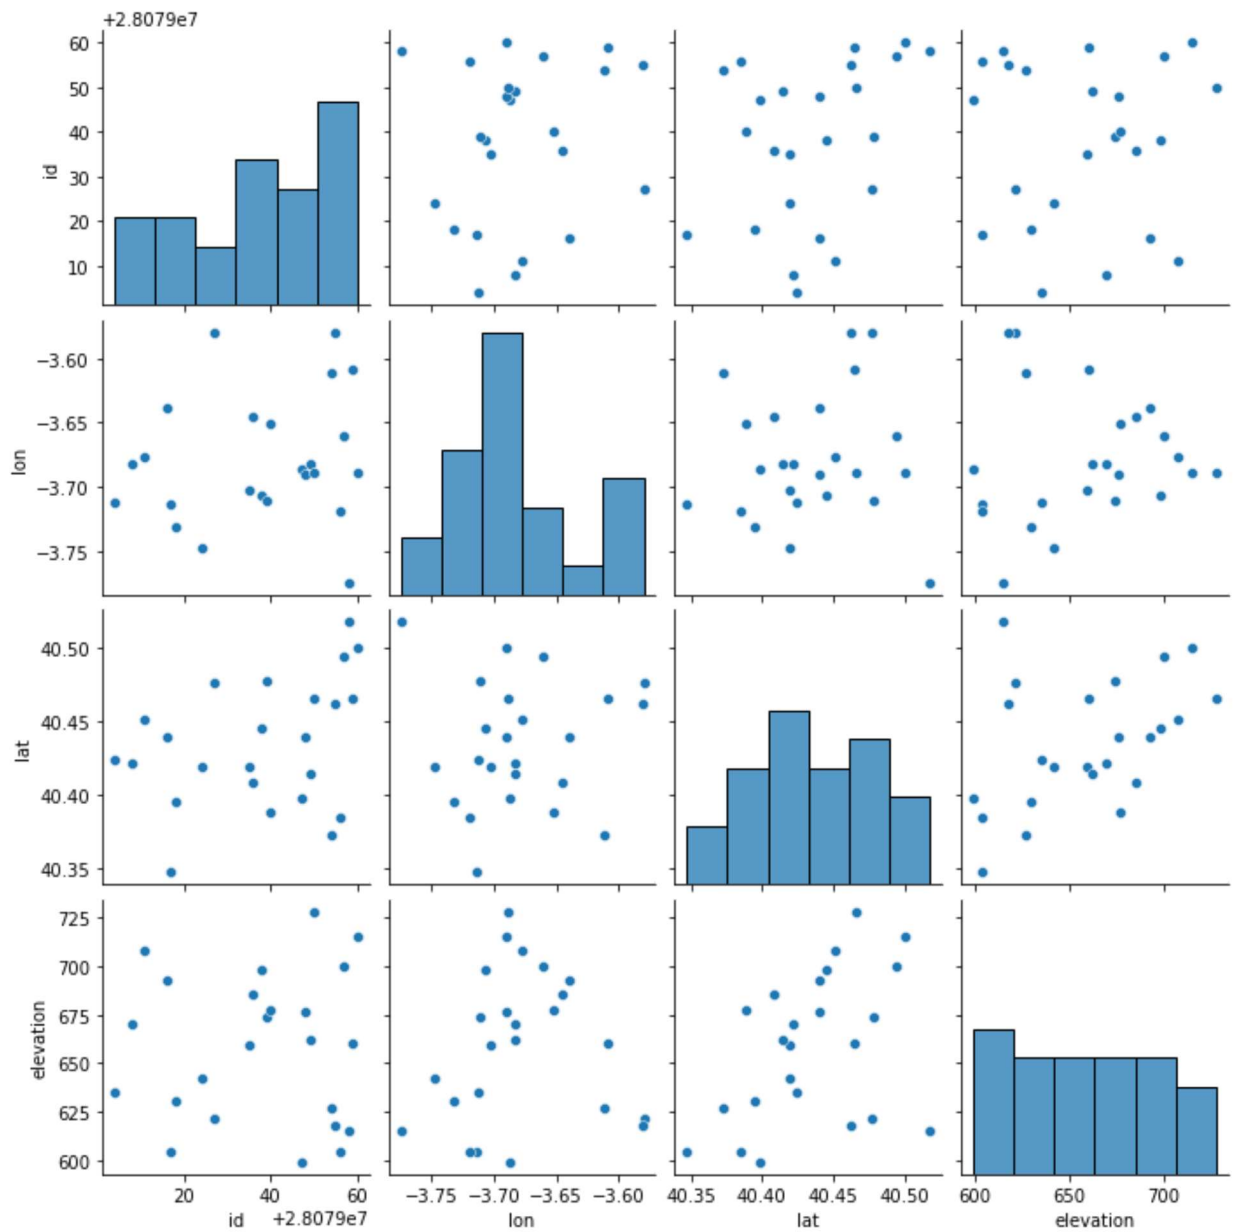
	id	lon	lat	elevation
<b>count</b>	2.400000e+01	24.000000	24.000000	24.000000
<b>mean</b>	2.807904e+07	-3.679019	40.434616	658.333333
<b>std</b>	1.799094e+01	0.049324	0.043022	38.295949
<b>min</b>	2.807900e+07	-3.774611	40.347139	599.000000
<b>25%</b>	2.807902e+07	-3.711718	40.405489	625.500000
<b>50%</b>	2.807904e+07	-3.687797	40.431875	661.000000
<b>75%</b>	2.807905e+07	-3.649968	40.465331	687.000000
<b>max</b>	2.807906e+07	-3.580031	40.518058	728.000000

In [20]: `df1=df[['id', 'name', 'address', 'lon', 'lat', 'elevation']]`

## EDA AND VISUALIZATION

```
In [21]: sns.pairplot(df1[0:50])
```


```
Out[21]: <seaborn.axisgrid.PairGrid at 0x1814e1f27f0>
```



```
In [24]: sns.distplot(df1['id'])
```


C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).  
warnings.warn(msg, FutureWarning)

```
Out[24]: <AxesSubplot:xlabel='id', ylabel='Density'>
```



```
In [25]: sns.heatmap(df1.corr())
```

```
Out[25]: <AxesSubplot:>
```



## TO TRAIN THE MODEL AND MODEL BUILDING

```
In [26]: x=df[['id']]
         y=df['elevation']
```

```
In [27]: from sklearn.model_selection import train_test_split
         x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

# Linear Regression

```
In [28]: from sklearn.linear_model import LinearRegression
lr=LinearRegression()
lr.fit(x_train,y_train)
```

Out[28]: LinearRegression()

```
In [29]: lr.intercept_
```

Out[29]: -10043641.837007152


```
In [30]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
coeff
```

Out[30]:

	Co-efficient
id	0.357715

```
In [31]: prediction =lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[31]: <matplotlib.collections.PathCollection at 0x18150ae9bb0>



## ACCURACY

```
In [32]: lr.score(x_test,y_test)
```

Out[32]: -0.4286690092552836

```
In [33]: lr.score(x_train,y_train)
```

Out[33]: 0.03214202174913028

## Ridge and Lasso

```
In [34]: from sklearn.linear_model import Ridge,Lasso
```

```
In [35]: rr=Ridge(alpha=10)  
rr.fit(x_train,y_train)
```

```
Out[35]: Ridge(alpha=10)
```

## Accuracy(Ridge)

```
In [36]: rr.score(x_test,y_test)
```

```
Out[36]: -0.4283479389212217
```

```
In [37]: rr.score(x_train,y_train)
```

```
Out[37]: 0.03214190424205121
```

```
In [38]: la=Lasso(alpha=10)  
la.fit(x_train,y_train)
```

```
Out[38]: Lasso(alpha=10)
```

```
In [39]: la.score(x_test,y_test)
```

```
Out[39]: -0.4144698759213483
```

## Accuracy(Lasso)

```
In [40]: la.score(x_train,y_train)
```

```
Out[40]: 0.0319060238163732
```

## Accuracy(Elastic Net)

```
In [41]: from sklearn.linear_model import ElasticNet  
en=ElasticNet()  
en.fit(x_train,y_train)
```

```
Out[41]: ElasticNet()
```

```
In [42]: en.coef_
```

```
Out[42]: array([0.35563738])
```

```
In [43]: en.intercept_
```

```
Out[43]: -9985303.982509833
```

```
In [44]: prediction=en.predict(x_test)
```

```
In [45]: en.score(x_test,y_test)
```

```
Out[45]: -0.42769432950583486
```

## Evaluation Metrics

```
In [46]: from sklearn import metrics
print(metrics.mean_absolute_error(y_test,prediction))
print(metrics.mean_squared_error(y_test,prediction))
print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
37.17529917322099
1939.8796702160532
44.0440650964015
```

## Logistic Regression

```
In [47]: from sklearn.linear_model import LogisticRegression
```

```
In [48]: feature_matrix=df[['id']]
target_vector=df['elevation']
```

```
In [49]: feature_matrix.shape
```

```
Out[49]: (24, 1)
```

```
In [50]: target_vector.shape
```

```
Out[50]: (24,)
```

```
In [51]: from sklearn.preprocessing import StandardScaler
```

```
In [52]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [53]: logr=LogisticRegression(max_iter=10000)
logr.fit(fs,target_vector)
```

```
Out[53]: LogisticRegression(max_iter=10000)
```

```
In [54]: observation=[[1]]
```

```
In [55]: prediction=logr.predict(observation)
print(prediction)
```

```
[604]
```

```
In [56]: logr.classes_
```

```
Out[56]: array([599, 604, 615, 618, 621, 627, 630, 635, 642, 659, 660, 662, 670,
        674, 676, 677, 685, 693, 698, 700, 708, 715, 728], dtype=int64)
```

```
In [57]: logr.score(fs,target_vector)
```

```
Out[57]: 0.16666666666666666
```

```
In [58]: logr.predict_proba(observation)[0][0]
```

```
Out[58]: 0.05149080255479361
```

```
In [59]: logr.predict_proba(observation)
```

```
Out[59]: array([[0.0514908 , 0.07628281, 0.06459573, 0.06111161, 0.02846647,
        0.05992926, 0.01982391, 0.00955813, 0.02542849, 0.03721458,
        0.06573175, 0.05391838, 0.01206135, 0.04186834, 0.05270507,
        0.04305319, 0.03836362, 0.0181069 , 0.04069126, 0.06344616,
        0.01416966, 0.06685293, 0.05512959]])
```

## Random Forest

```
In [60]: from sklearn.ensemble import RandomForestClassifier
```

```
In [61]: rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[61]: RandomForestClassifier()
```

```
In [62]: parameters={'max_depth':[1,2,3,4,5],
        'min_samples_leaf':[5,10,15,20,25],
        'n_estimators':[10,20,30,40,50]
}
```

```
In [63]: from sklearn.model_selection import GridSearchCV
grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")
grid_search.fit(x_train,y_train)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\model\_selection\\_split.py:666: UserWarning: The least populated class in y has only 1 members, which is less than n\_splits=2.

warnings.warn(("The least populated class in y has only %d"

```
Out[63]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
        param_grid={'max_depth': [1, 2, 3, 4, 5],
        'min_samples_leaf': [5, 10, 15, 20, 25],
        'n_estimators': [10, 20, 30, 40, 50]},
        scoring='accuracy')
```

```
In [64]: grid_search.best_score_
```

```
Out[64]: 0.125
```

```
In [65]: rfc_best=grid_search.best_estimator_
```

```
In [66]: from sklearn.tree import plot_tree

plt.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['a','b','c'],
```

```
Out[66]: [Text(2232.0, 1087.2, 'gini = 0.875\nsamples = 9\nvalue = [3, 0, 1, 2, 2, 0, 0, 2, 0, 2, 2, 1, 0, 1, 0, 2, 2, 1, 0, 1\n0]\n\nclass = a')]
```

```
gini = 0.875
samples = 9
value = [3, 0, 1, 2, 2, 0, 0, 2, 0, 2, 2, 1, 0, 1, 0, 2, 2, 1, 0, 1
0]
class = a
```

## Conclusion

### Accuracy

**Linear Regression:0.03214202174913028**

**Ridge Regression:0.03214190424205121**

**Lasso Regression:0.0319060238163732**

**ElasticNet Regression:-0.42769432950583486**

**Logistic Regression:0.16666666666666666**

**Random Forest:0.125**

**Logistic Regression is suitable for this dataset**



